

Llama.Lisp: AI First Compiler Framework

Rationale and Design

Sasank Chilamkurthy



Outline

- Why my own compiler framework
- Insides of a compiler
- Principles of my design

This work wouldn't have been possible without

- **Vedanth Padmaraman**
- **Adithya L Bhat**



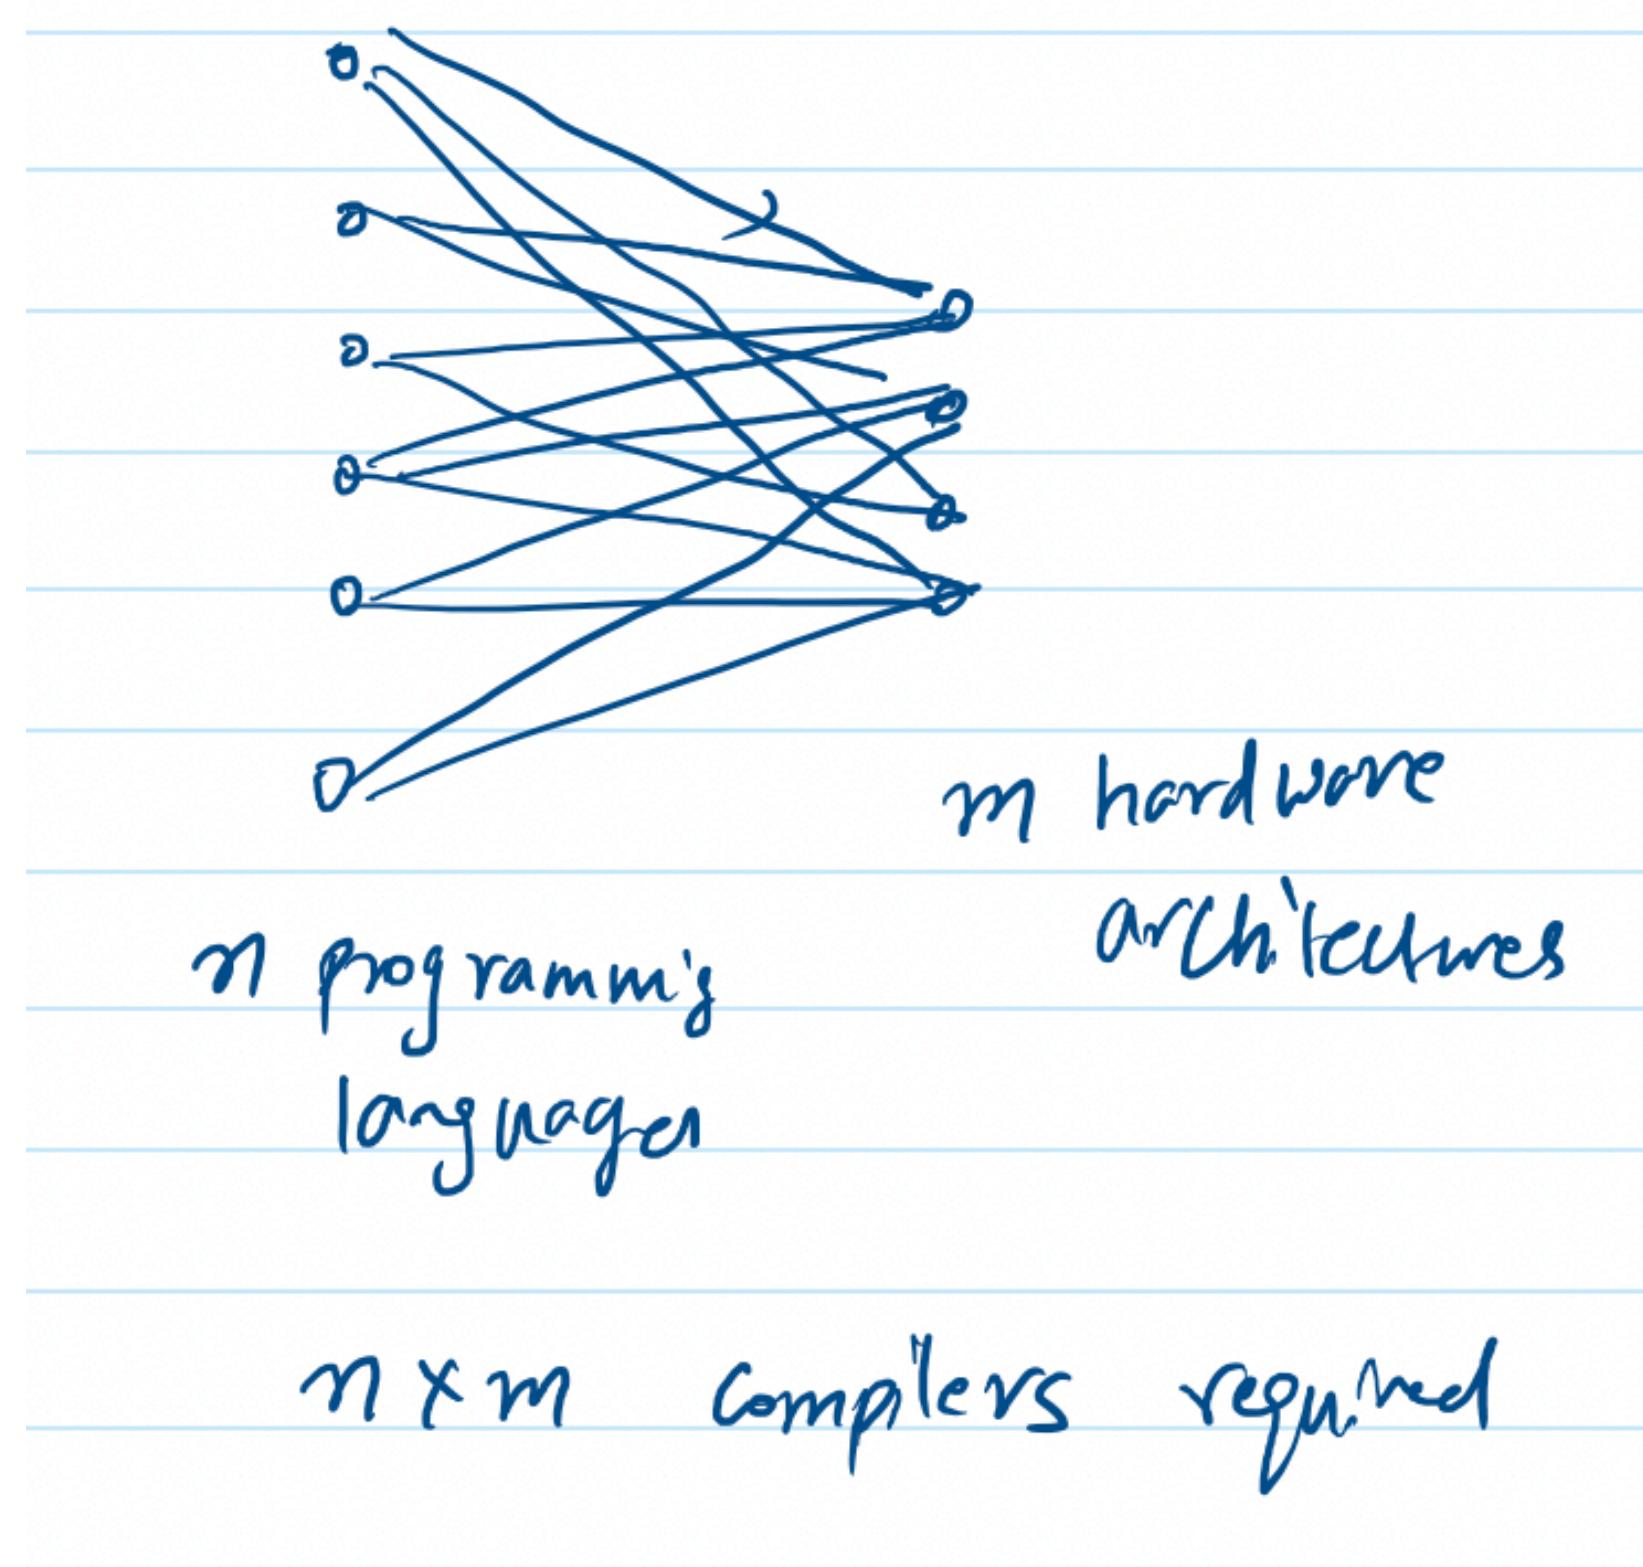
Why My Own Compiler Framework?



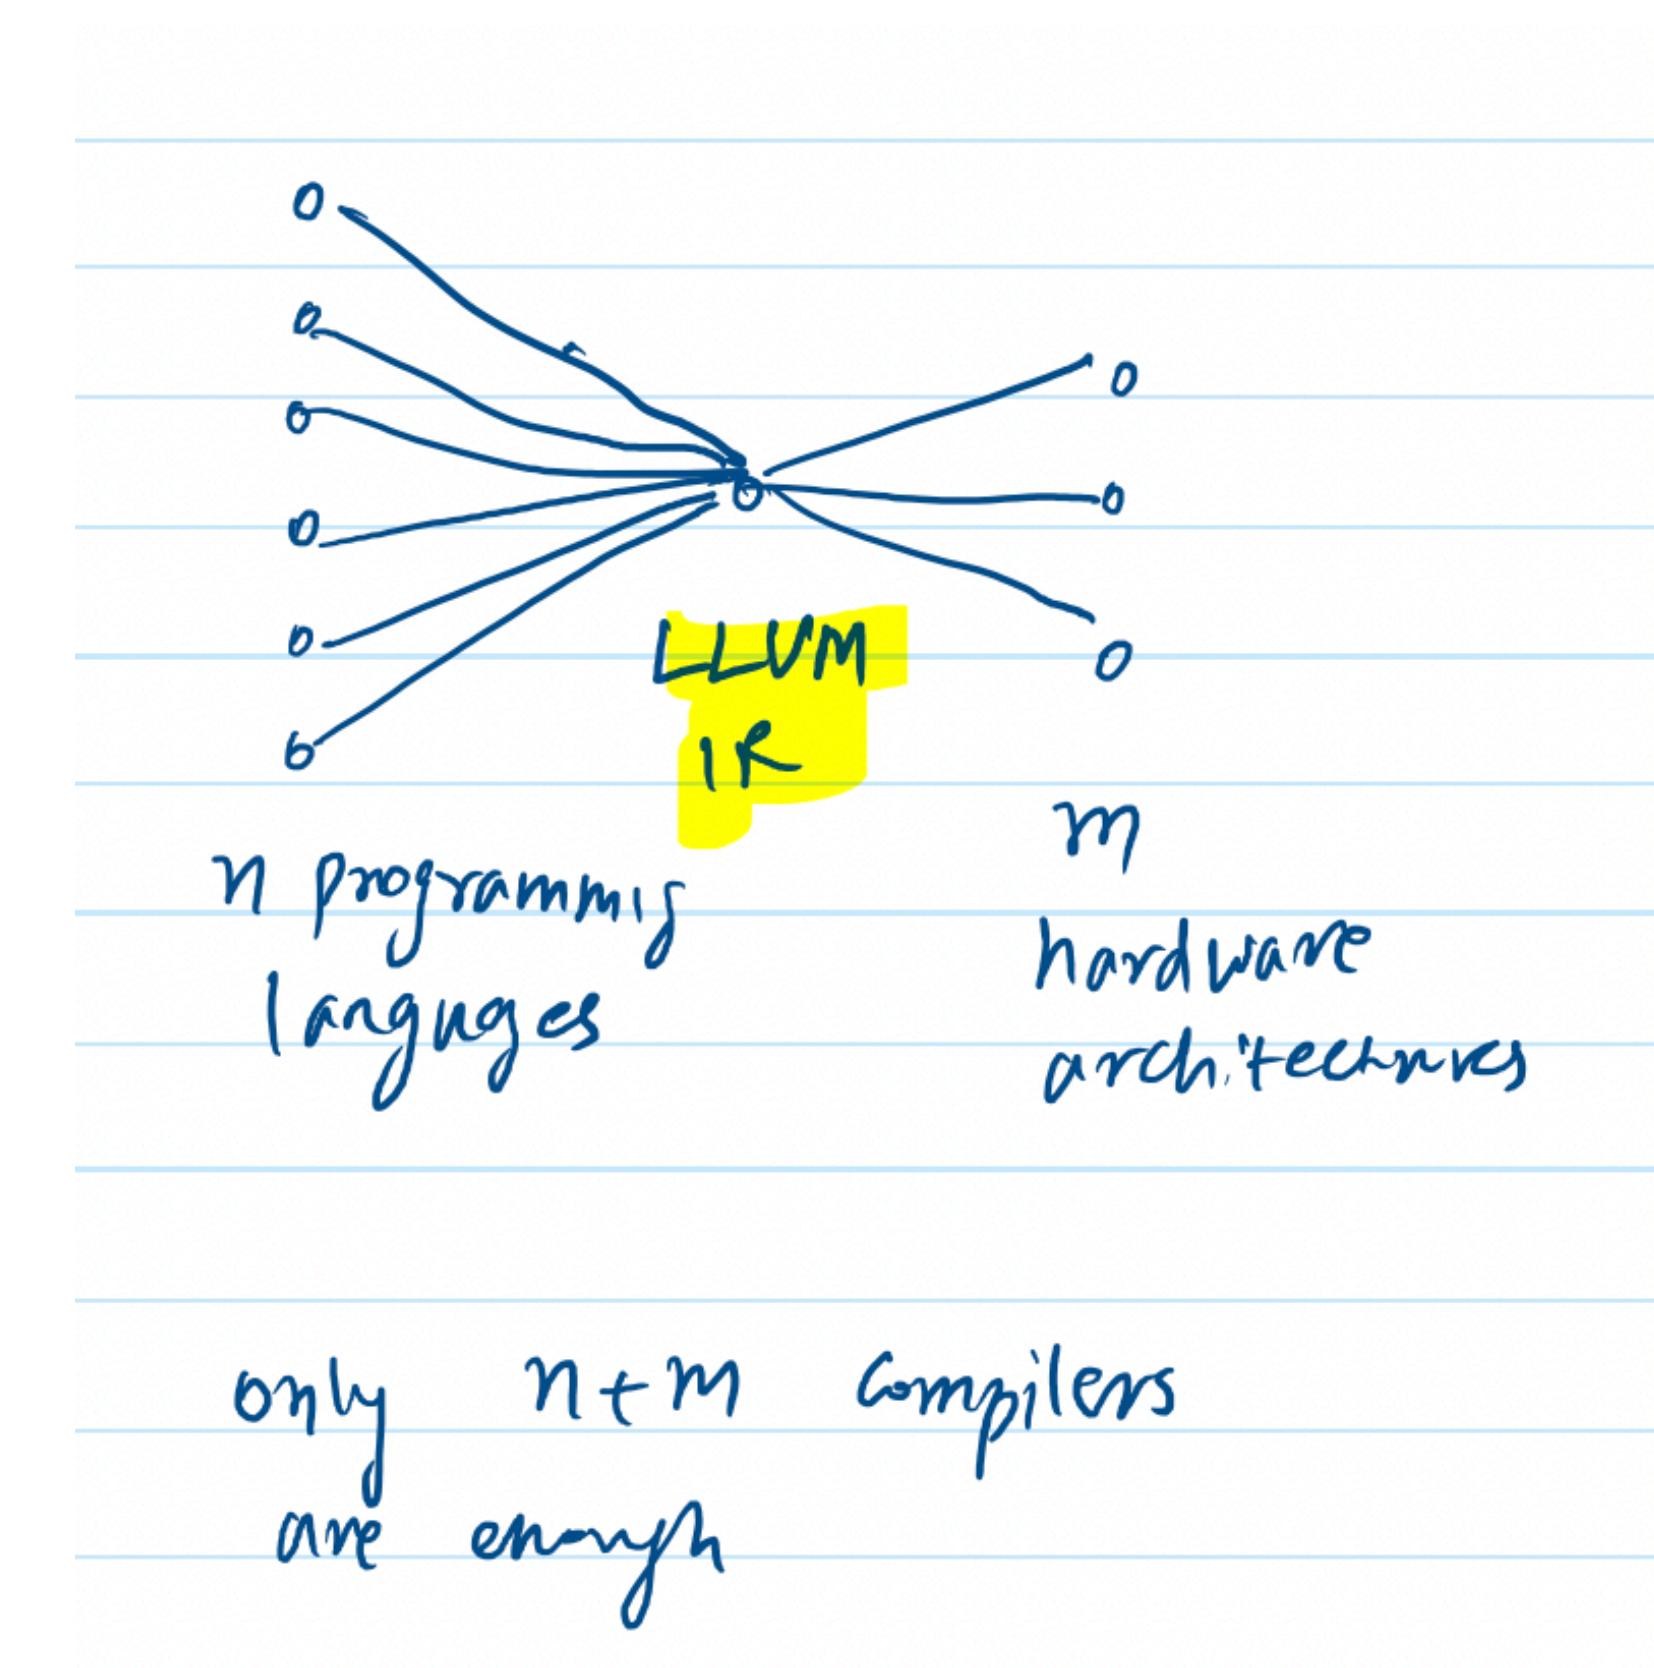
Compilers are Translators

Just like English to Kannada but Rust to x86

Without IR



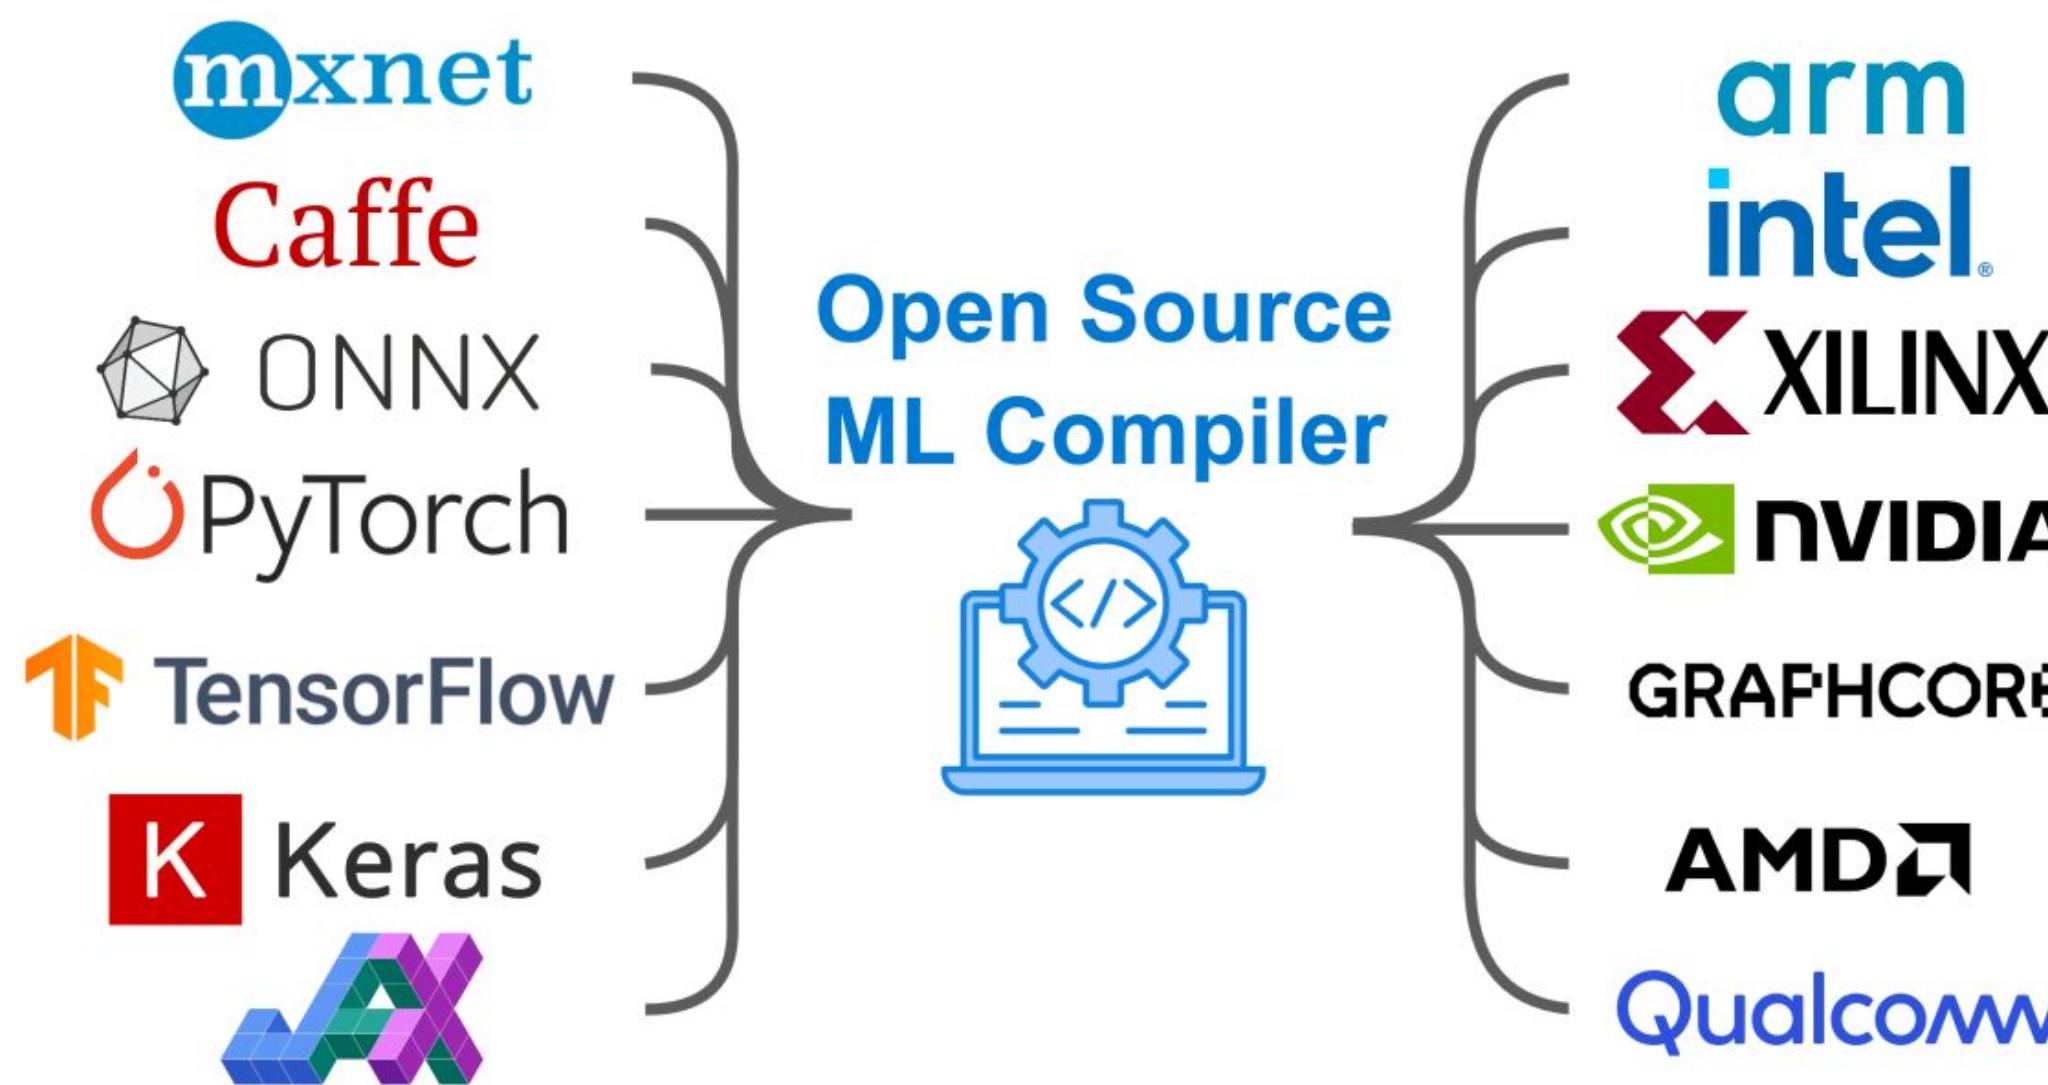
With IR



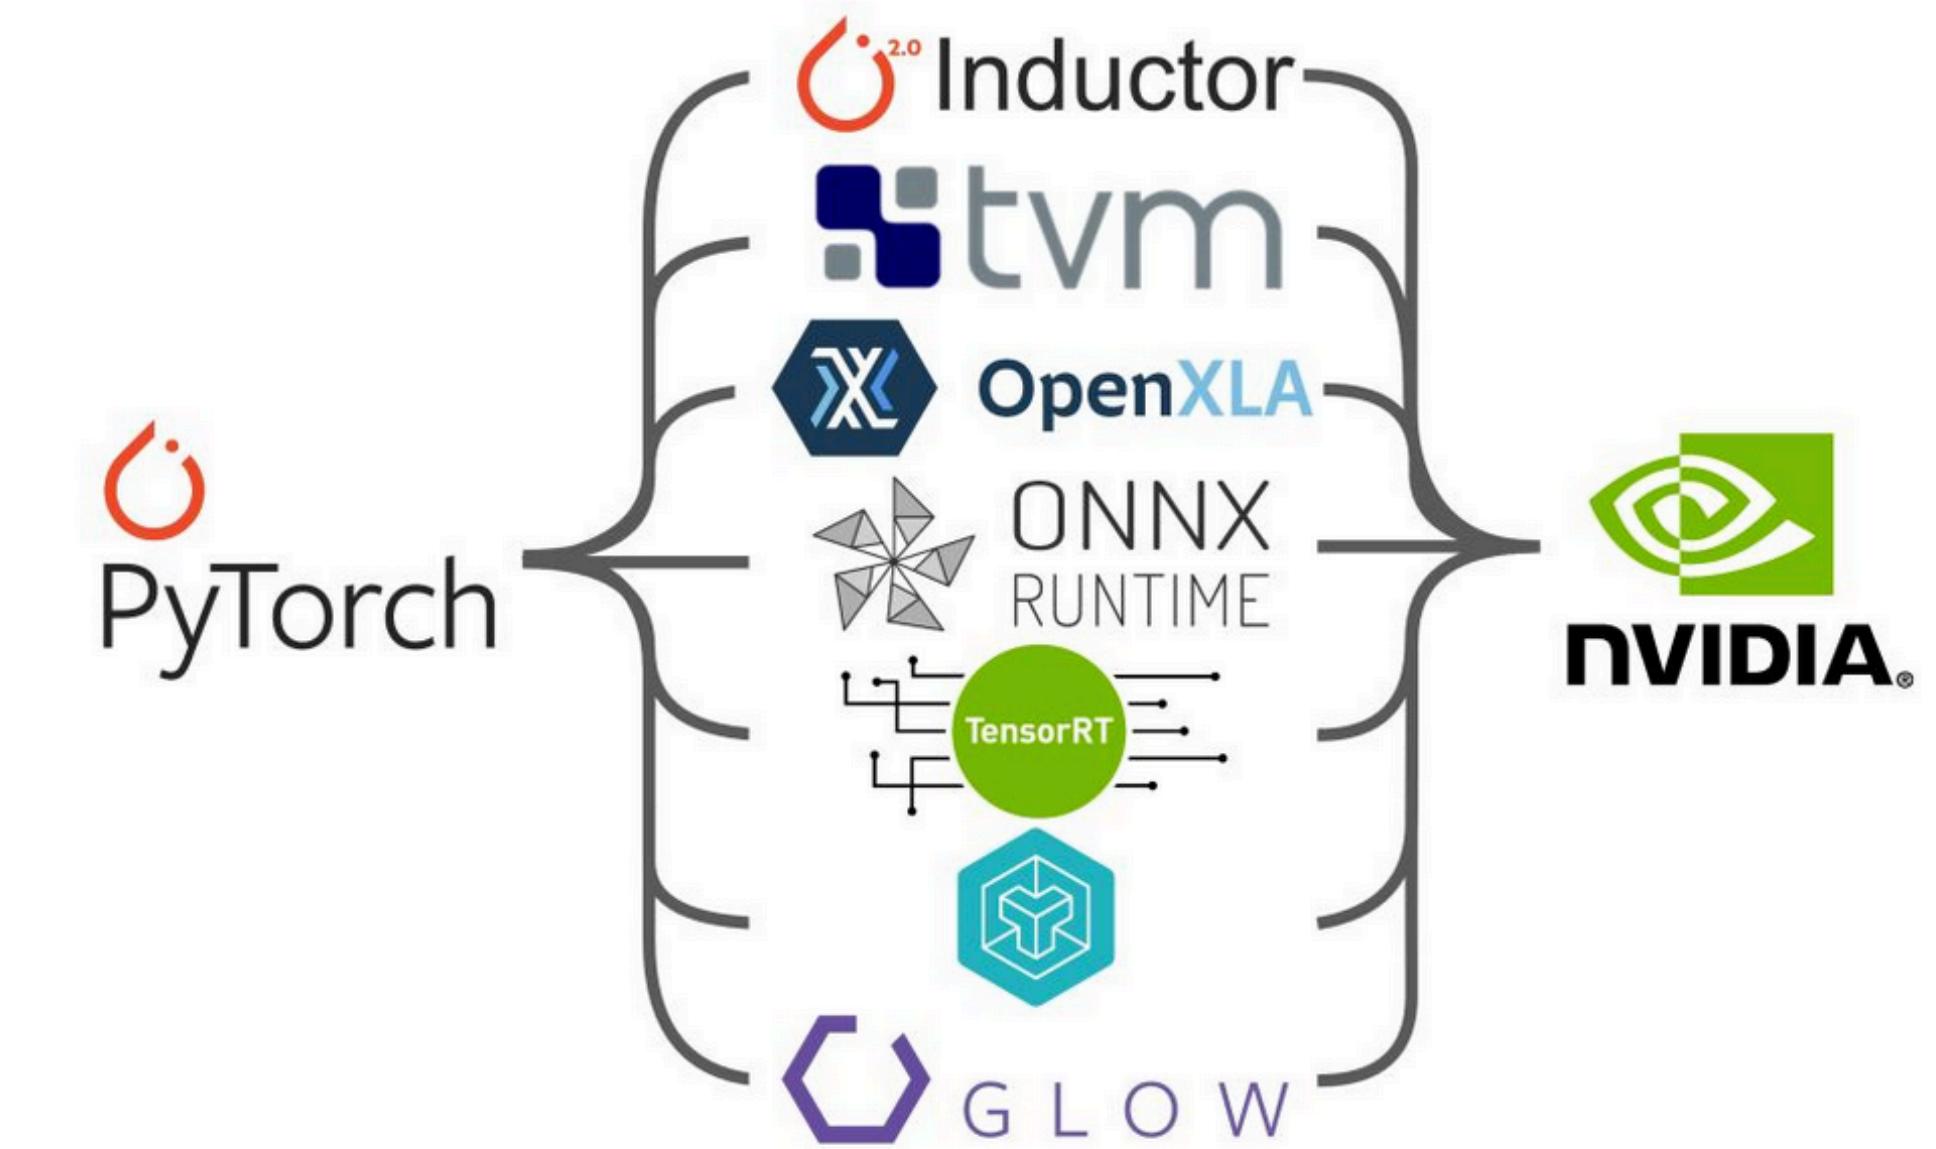


AI Compilers: Expectation vs Reality

Expectation



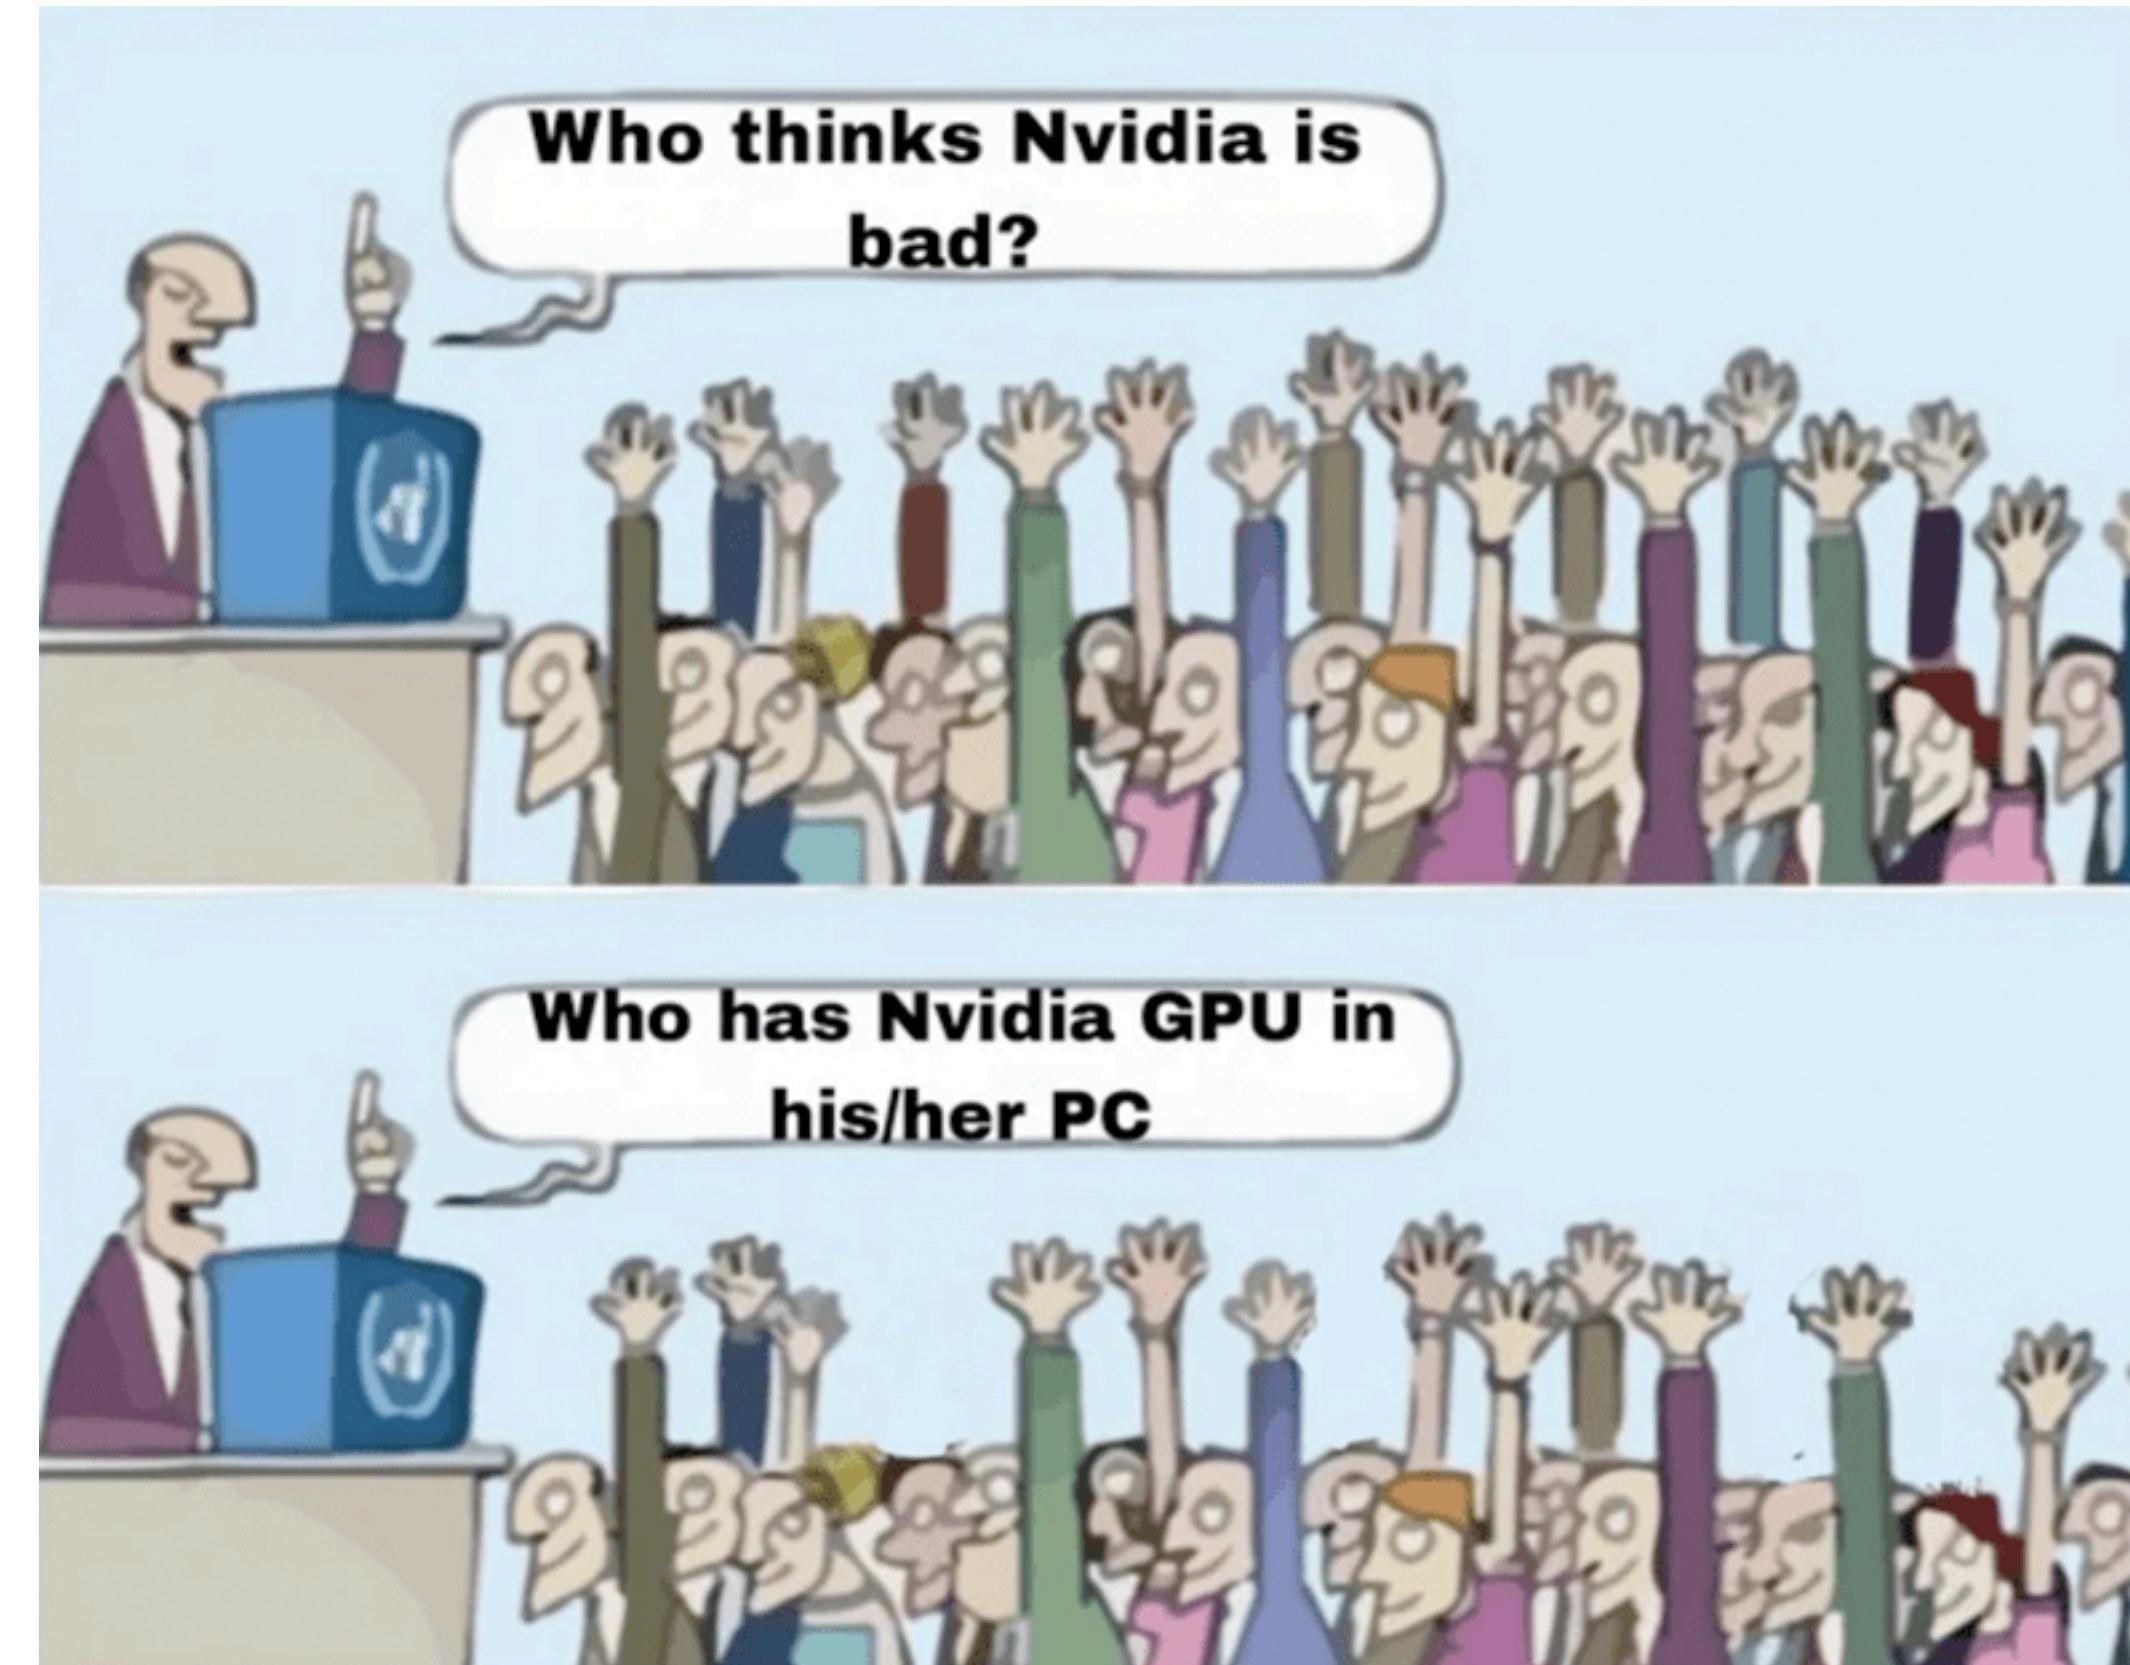
Reality





The Consequence

We're Stuck with Single Vendor



https://www.reddit.com/r/pcmasterrace/comments/14w6ar4/sad_but_true/



GPU Poor

GPUs unaffordable by most people



https://x.com/var_epsilon/status/1712301543980925331/photo/1

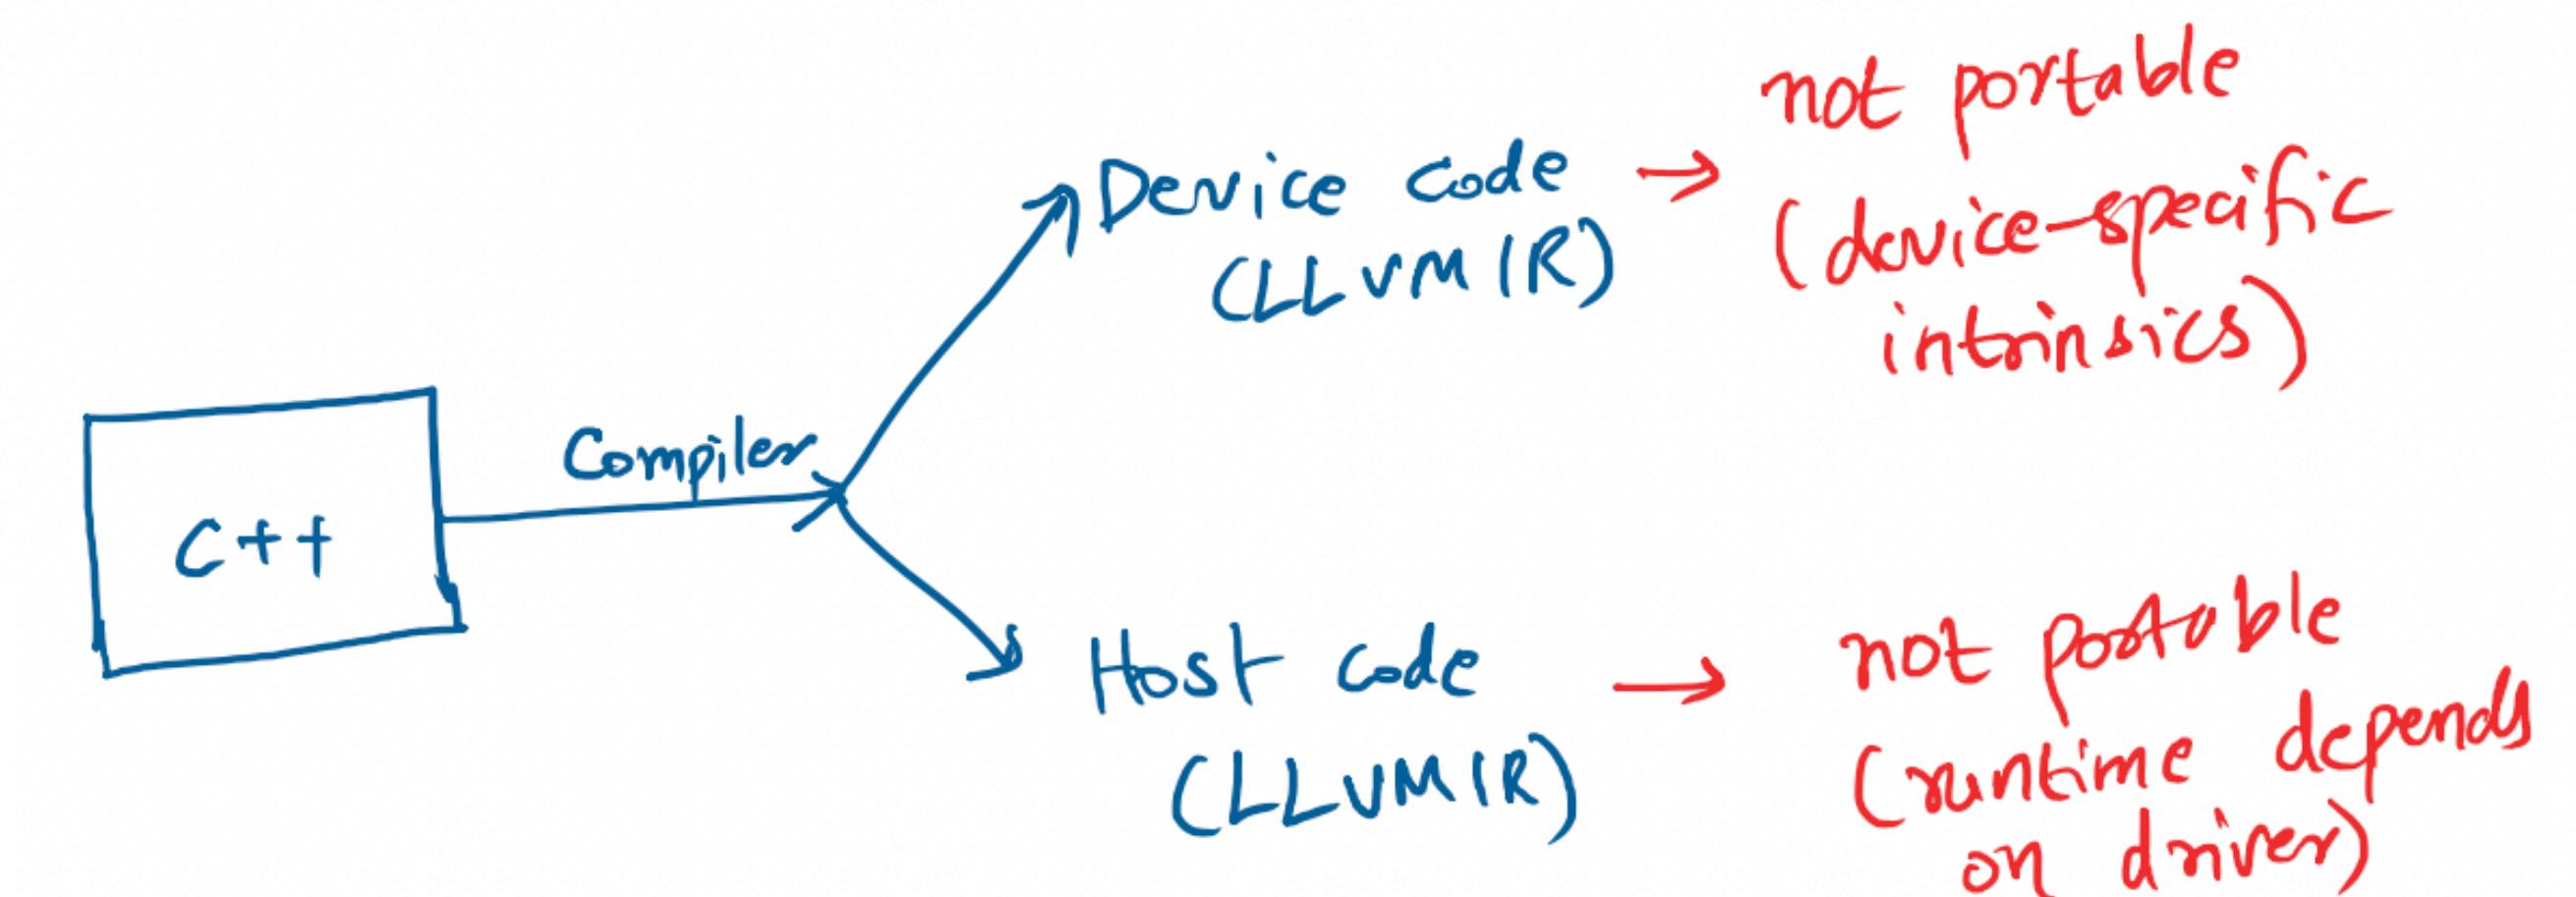
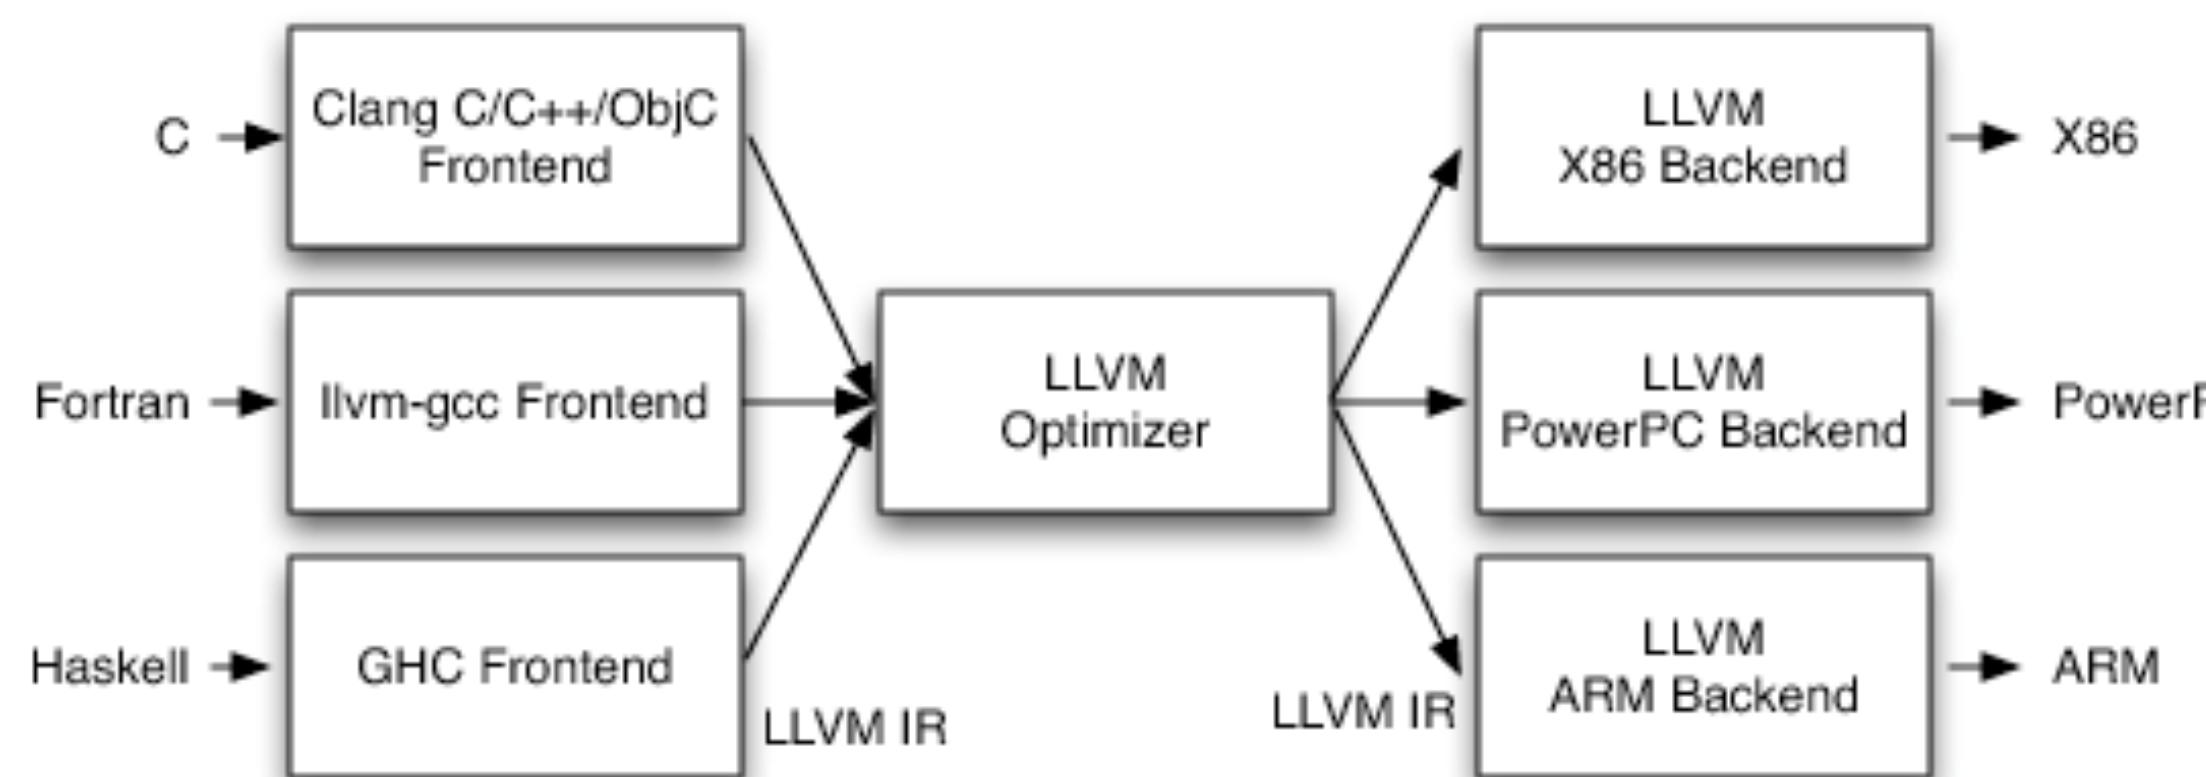
But, We have LLVM?

Unfortunately its design doesn't work for GPUs

[Home](#) • [My Writing](#) • [Classic Papers](#) • [About Me](#) • [Toggle Dark Mode](#)

Intermediate Representations for GPUs: LLVM Does Not Cut it

Sasank Chilamkurthy | 05 April 2024 | 11 minutes to read.





Introducing Llama.lisp

My Own Compiler Framework

```
1 (define ((n-squares void) (n int))
2   ; Declare `i` as an integer
3   (declare (i int))
4   ; Iterate over the first n integers
5   (for ((set i 0)
6         (lt i n)
7         (set i (add i 1)))
8     ; `sq` is scoped to the `for` loop body
9     (declare (sq int))
10    (set sq (mul i i))
11    (call print sq))
12
13  (call print sq) ; Error: `sq` is undeclared
14  (ret))
```

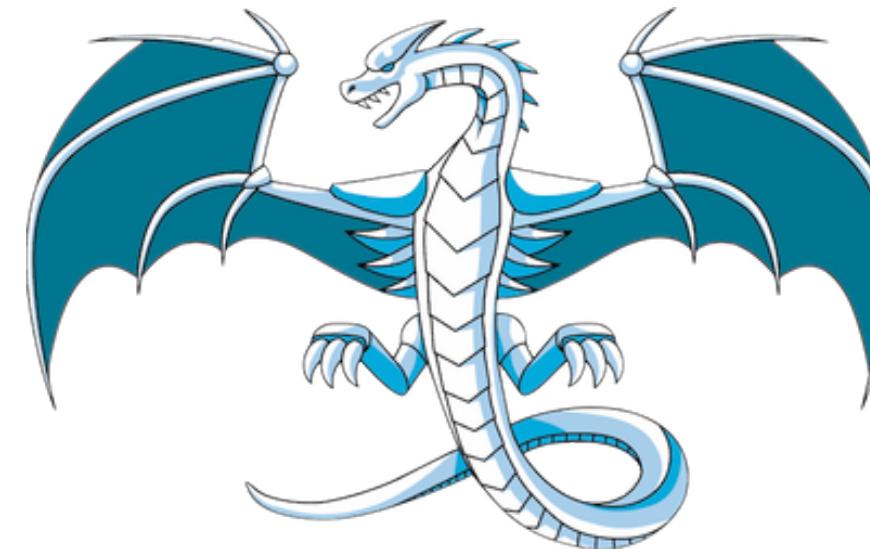
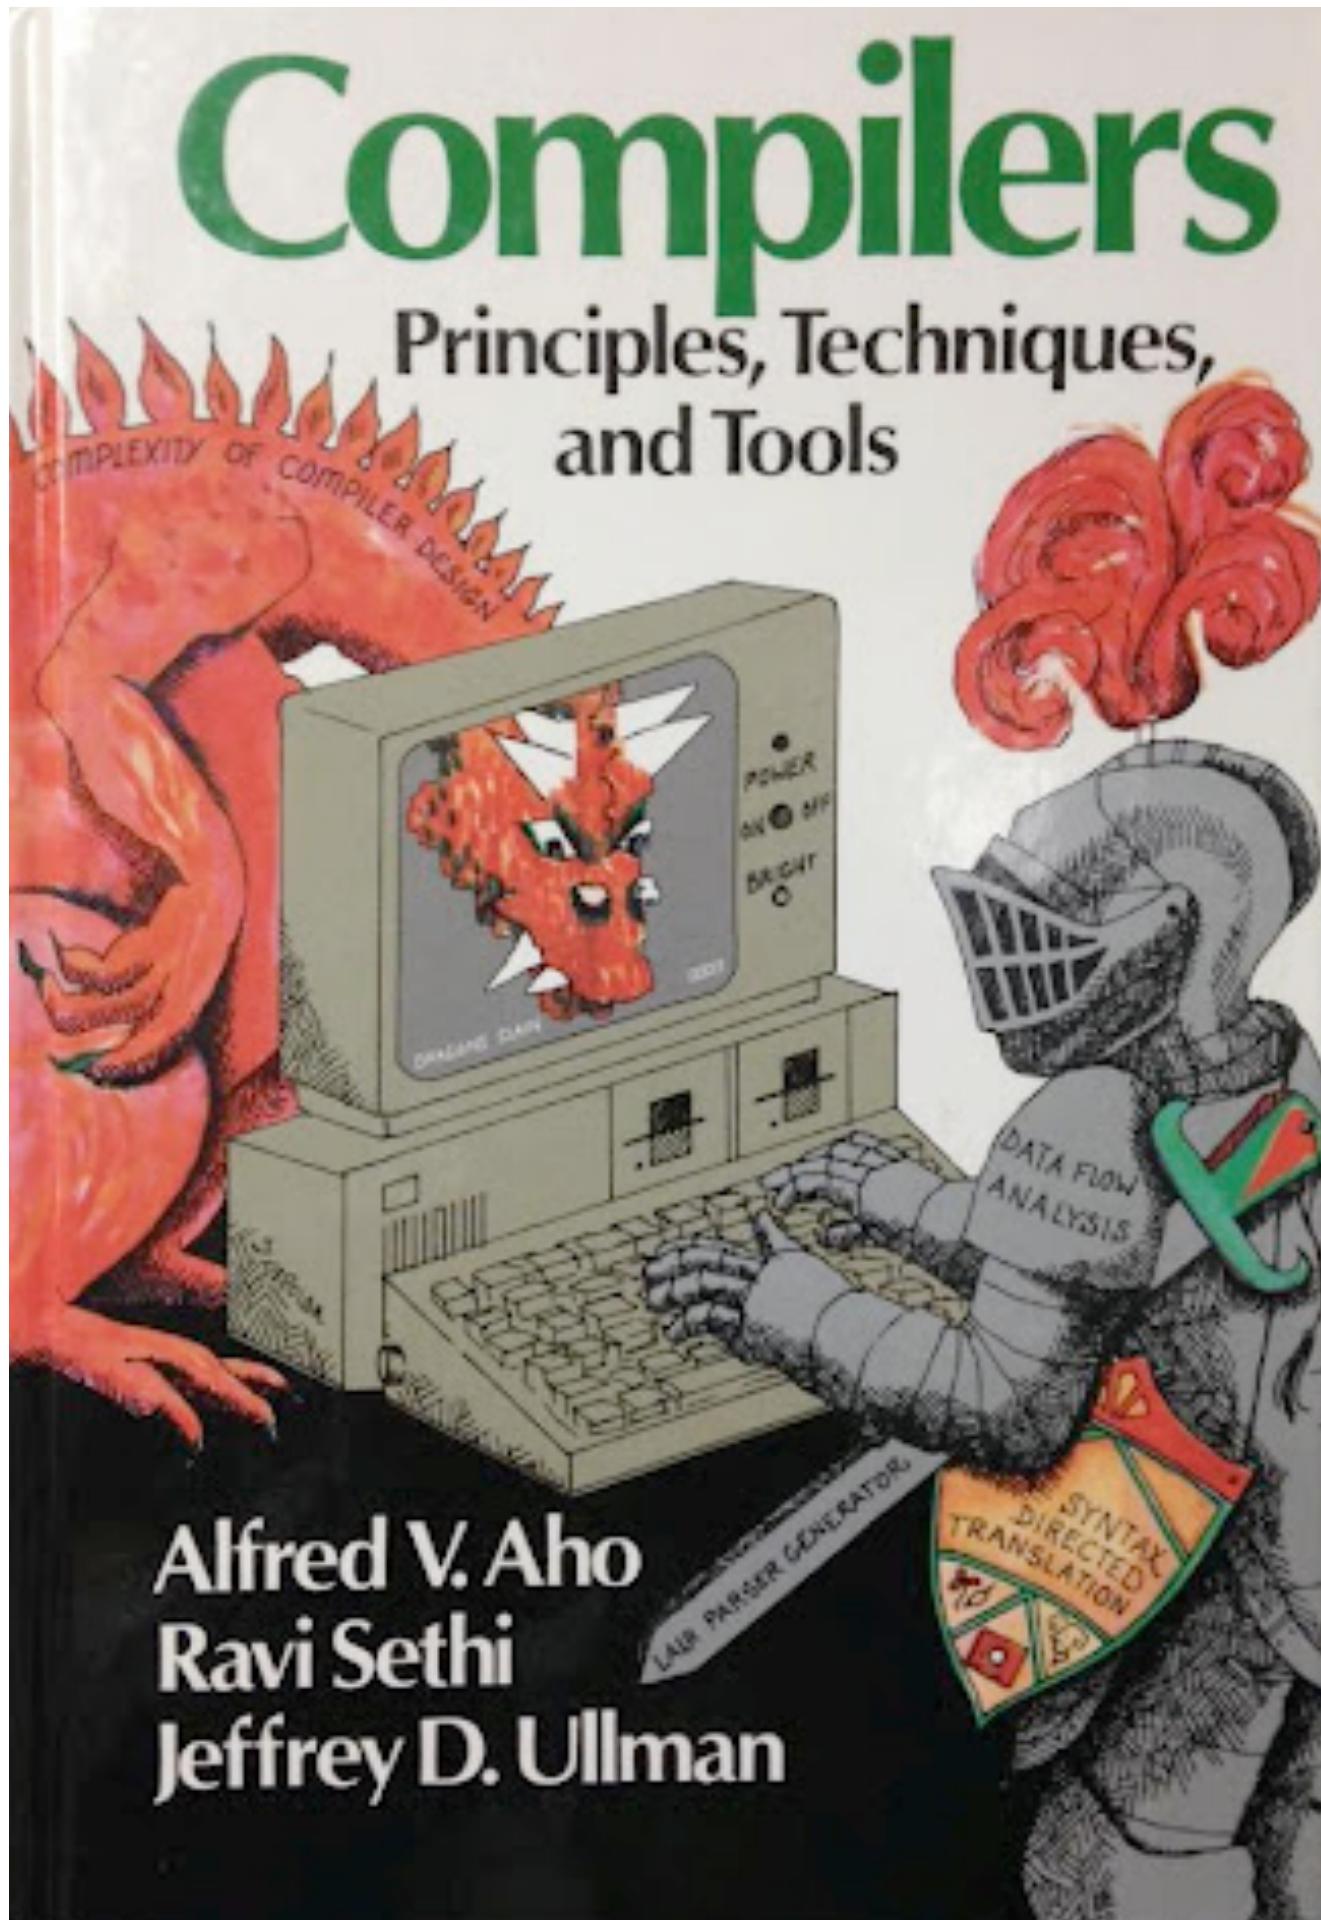
```
1 (bril-define ((add5 int) (n int))
2   (set (five int) (const 5))
3   (set (sum int) (add n five))
4   (ret sum))
5
6 (bril-define ((main int)))
7   (set (a int) (const 9))
8   (set (b int) (call add5 a))
9   (ret b))
```



Insides of a Compiler



Compilers are Dragons



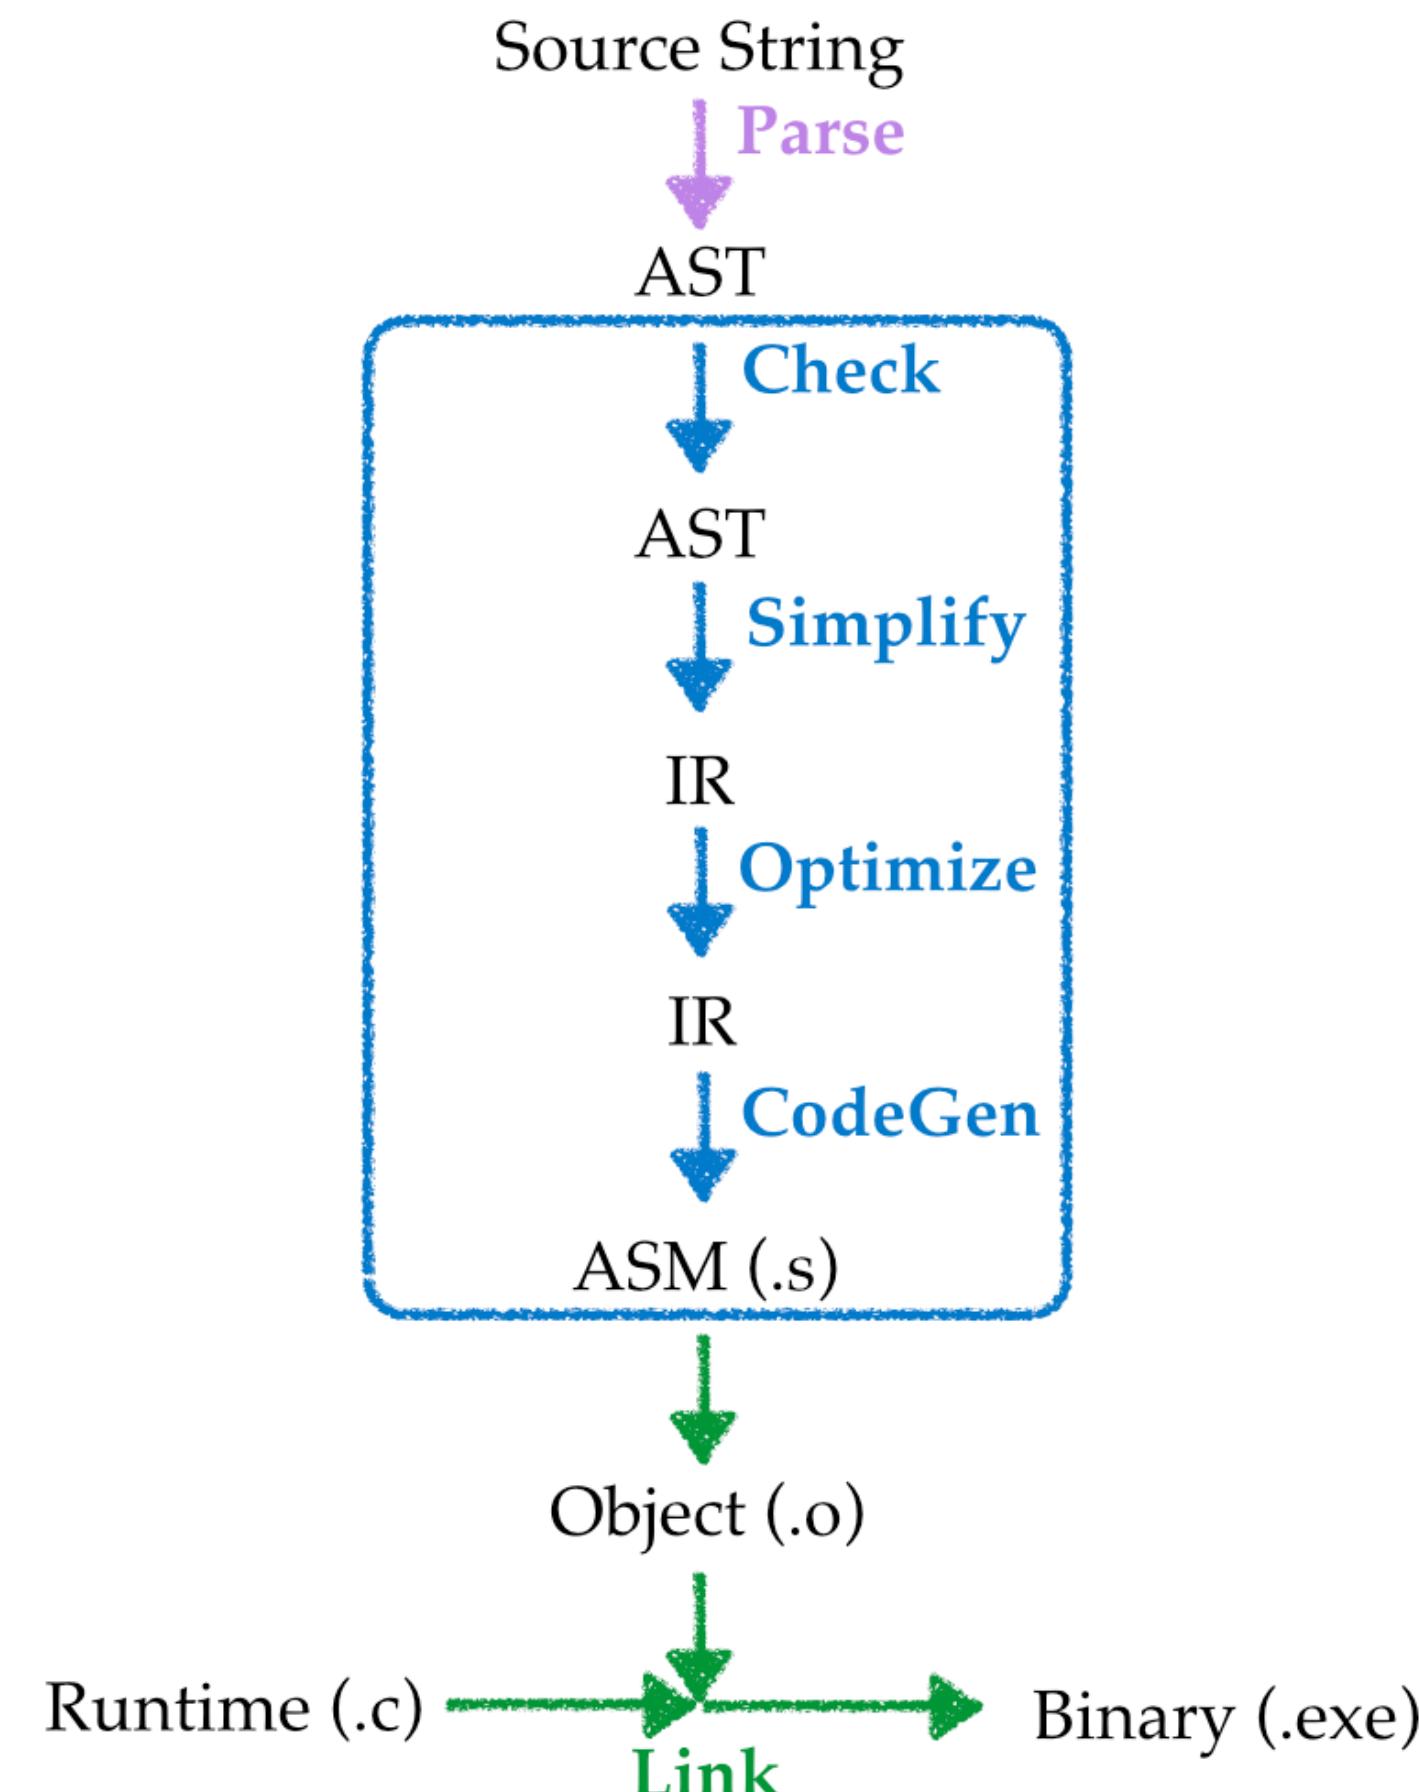
Not a pokemon lol





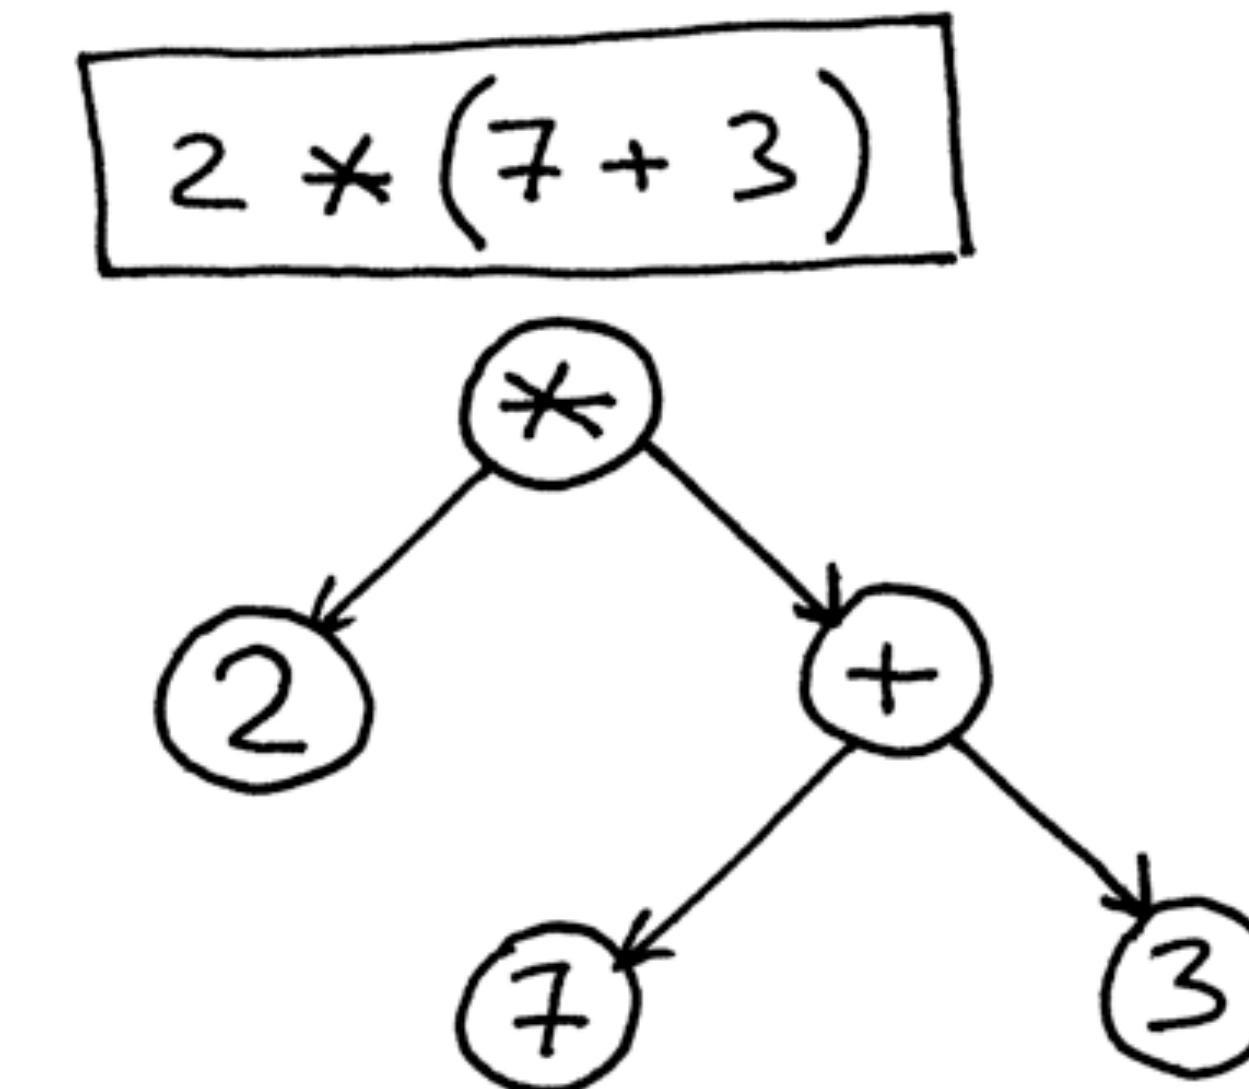
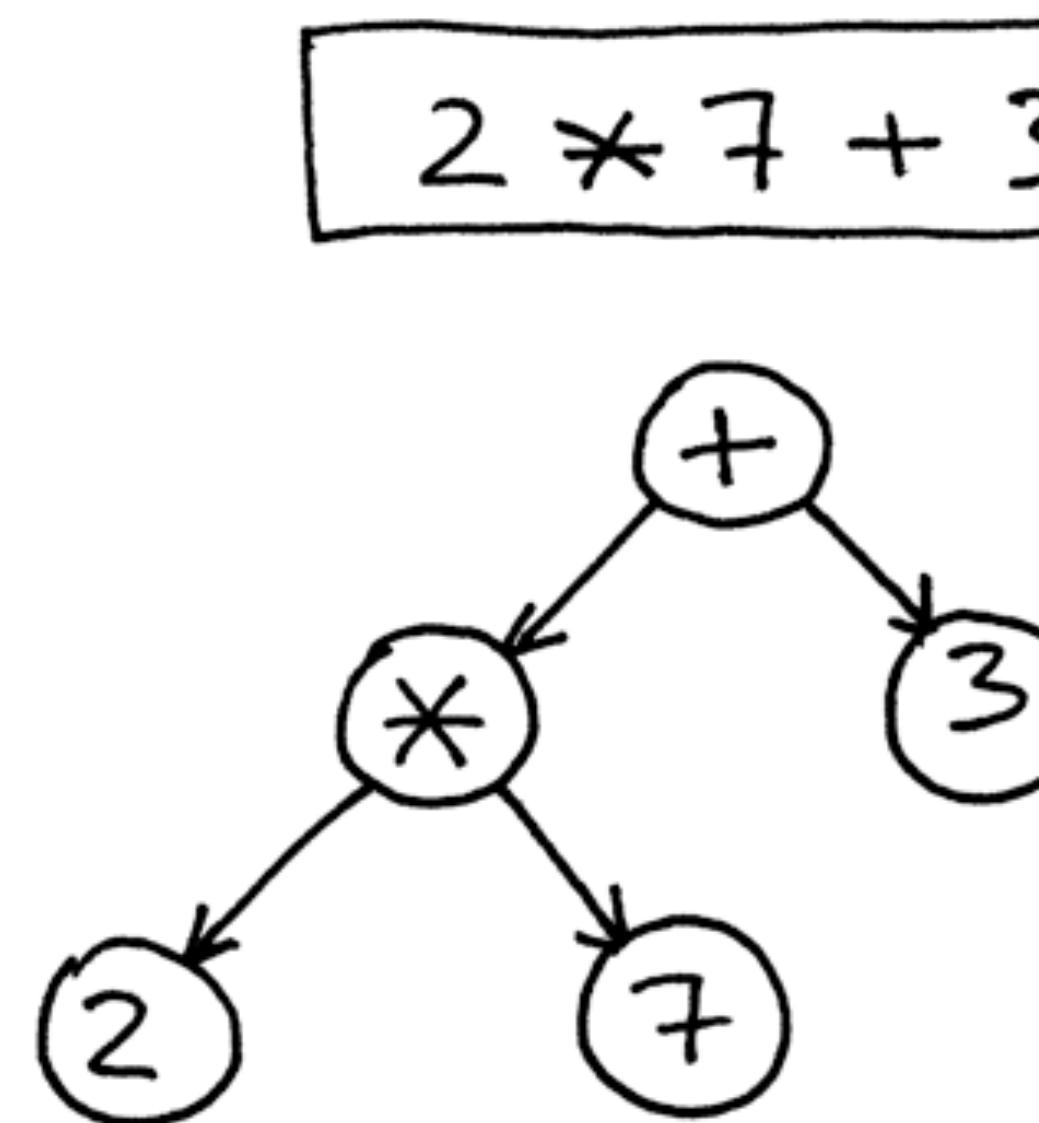
Parts of a Compiler

Innards of the dragon



Parser

Converts Program to a Abstract Syntax Tree





Intermediate Representation

Goes closer to assembly

```
int max(int a, int b) {  
    if (a > b) {  
        return a;  
    } else {  
        return b;  
    }  
}
```

```
define i32 @max(i32 %a, i32 %b) {  
entry:  
    %retval = alloca i32, align 4  
    %0 = icmp sgt i32 %a, %b  
    br i1 %0, label %btrue, label %bfalse  
  
btrue:                                ; preds = %2  
    store i32 %a, i32* %retval, align 4  
    br label %end  
  
bfalse:                                ; preds = %2  
    store i32 %b, i32* %retval, align 4  
    br label %end  
  
end:                                     ; preds = %btrue, %bfalse  
    %1 = load i32, i32* %retval, align 4  
    ret i32 %1  
}
```



Code Generation

IR is transformed to Assembly

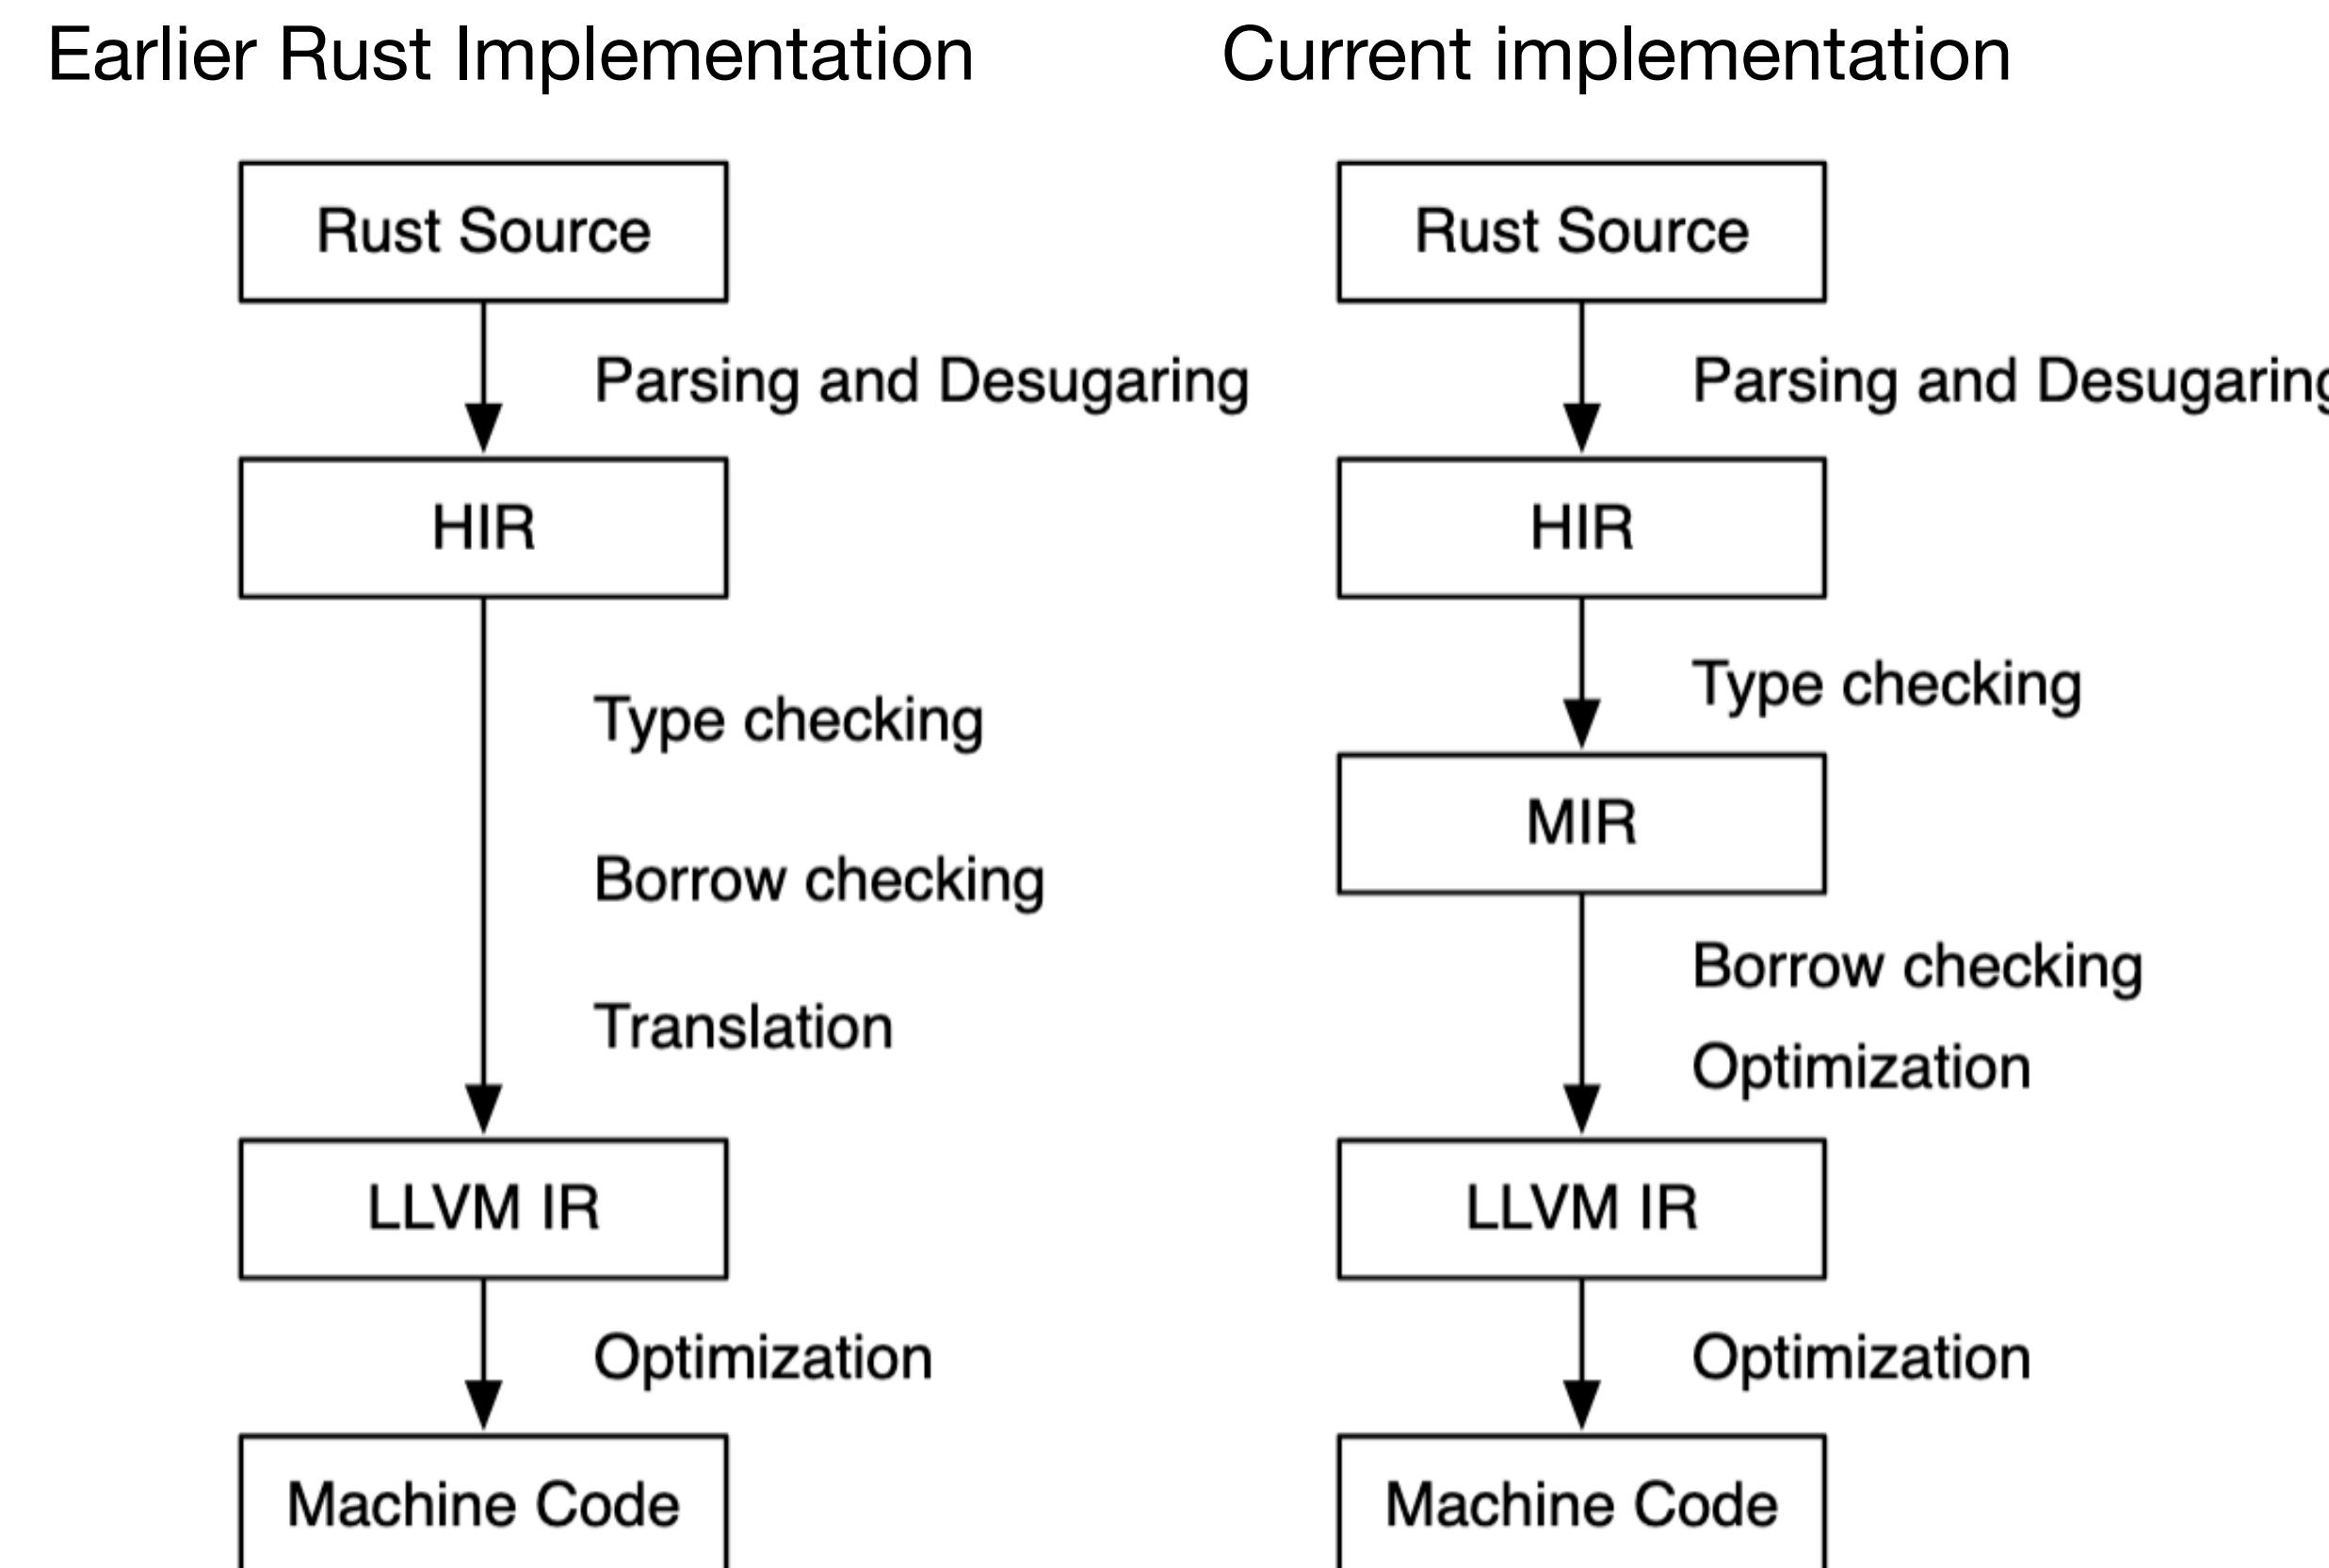
```
1 ; ModuleID = ""
2 target triple = "unknown-unknown-unknown"
3 target datalayout = ""
4
5 define i32 @"add5"(i32 %"n")
6 {
7 entry:
8     %"sum" = add i32 %"n", 5
9     ret i32 %"sum"
10 }
11
12 define i32 @"main"()
13 {
14 entry:
15     %"b" = call i32 @"add5"(i32 9)
16     ret i32 %"b"
17 }
```

```
.text
.file  "x.ll"
.globl add5
.p2align 4, 0x90
.type   add5,@function
add5:                                # @add5
    .cfi_startproc
# %bb.0:                                # %entry
    movl  4(%esp), %eax
    addl  $5, %eax
    retl
.Lfunc_end0:
    .size  add5, .Lfunc_end0-add5
    .cfi_endproc
                                # -- End function
                                # -- Begin function main
.globl main
.p2align 4, 0x90
.type   main,@function
main:                                # @main
    .cfi_startproc
# %bb.0:                                # %entry
    pushl $9
    .cfi_adjust_cfa_offset 4
    calll add5@PLT
    addl  $4, %esp
    .cfi_adjust_cfa_offset -4
    retl
.Lfunc_end1:
    .size  main, .Lfunc_end1-main
    .cfi_endproc
                                # -- End function
.section ".note.GNU-stack","",@progbits
```



Multiple Intermediate Representations

To tame the complexity





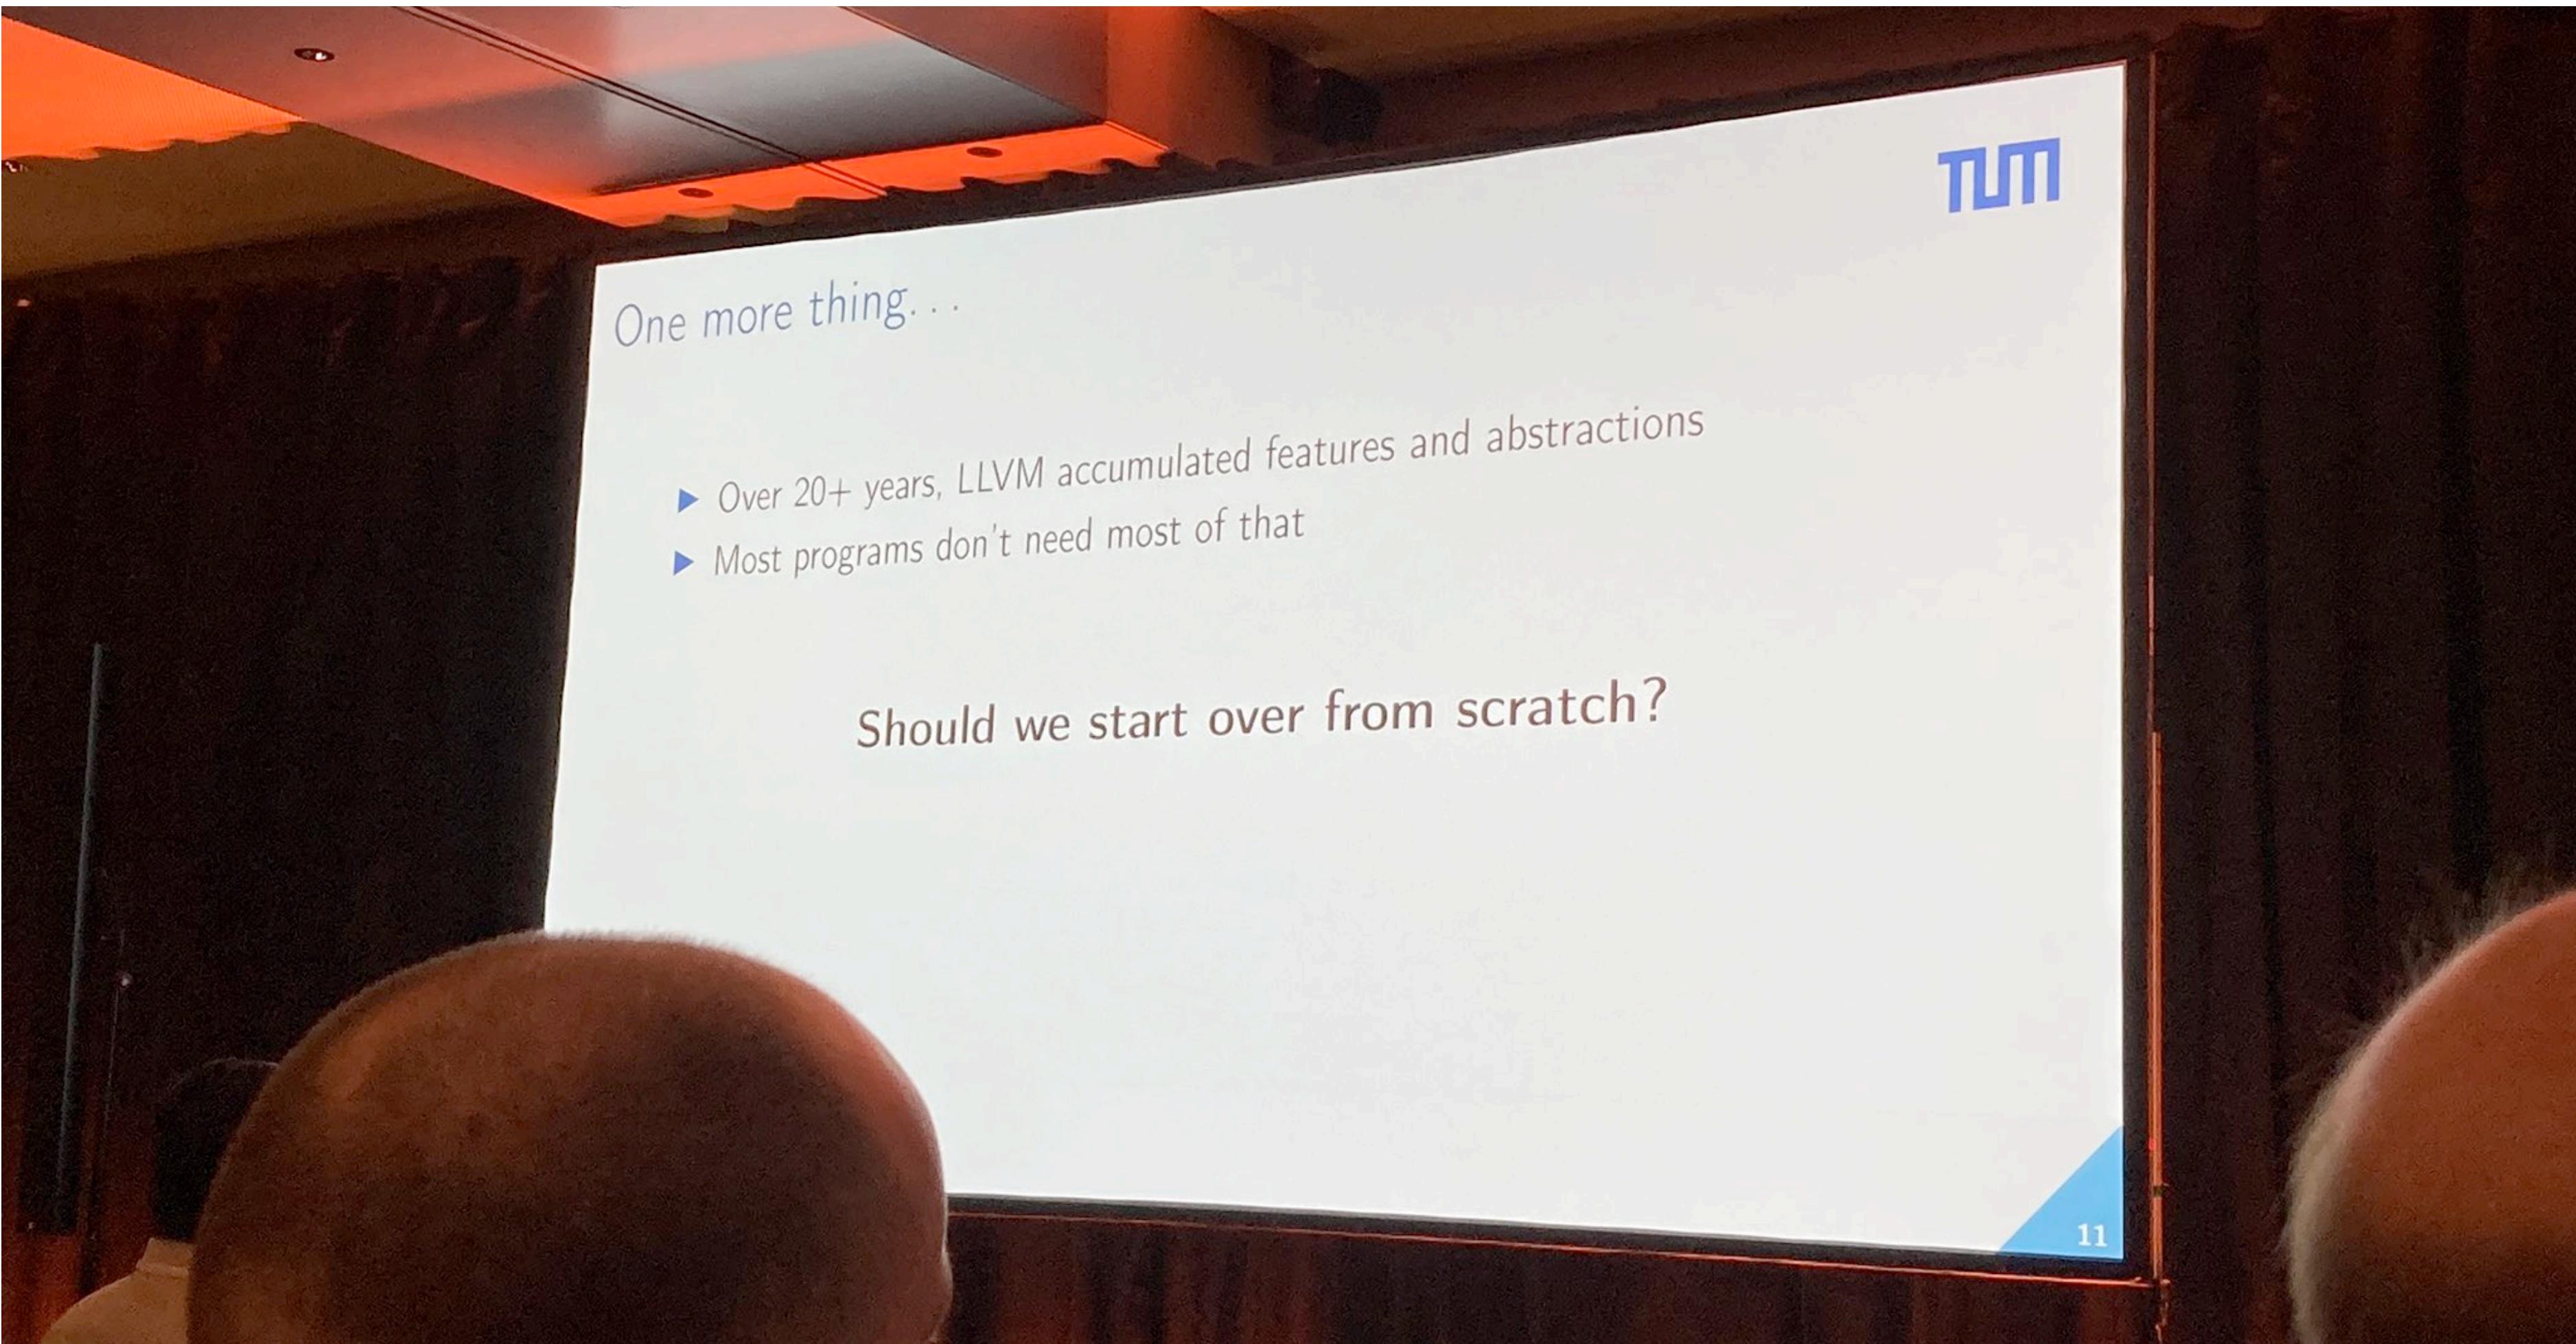
Design of llama.lisp



Philosophy

- Bring down the barrier of compiler writing significantly
- Avoid parsing, focus on semantics: JSON as IR
- Use the best tools for a job: implement in any language
- Highly extensible: Layers of languages
- Performance of generated code important

LLVM became too complex





BRIL: Simplification of LLVM IR

Lessons

Lesson 1: Welcome & Overview

[discussion](#) [video](#) [tasks](#) due September 9

Lesson 2: Representing Programs

[discussion](#) [representing_programs](#) [getting_started_with_BrIL](#)
[tasks](#) due September 14

Lesson 3: Local Analysis & Optimization

[discussion](#) [simple_dead_code_elimination](#) [local_value_numbering](#)
[tasks](#) due September 21

Lesson 4: Data Flow

[discussion](#) [data_flow](#) [implementation_task](#)
[tasks](#) due September 28

Lesson 5: Global Analysis & SSA

[discussion](#) [global_analysis_optimization](#) [static_single_assignment](#)
[tasks](#) due October 7

Lesson 6: LLVM

[discussion](#) [introduction_to_LLVM](#) [writing_an_LLVM_pass](#)
[tasks](#) due October 19

Lesson 7: Loop Optimization

[discussion](#) [video](#) [tasks](#) due November 2

Lesson 8: Interprocedural Analysis

[discussion](#) [video](#)

Lesson 9: Alias Analysis

[discussion](#) [video](#)

Lesson 10: Memory Management

[discussion](#) [video](#) [tasks](#) due November 9

Lesson 11: Dynamic Compilers

[discussion](#) [Dynamic_Compilers](#) [Tracing_via_Speculation](#)
[tasks](#) due December 7

Lesson 12: Program Synthesis

Bril: A Compiler Intermediate Representation for Learning

- Instruction Oriented
- Minimal and ruthlessly regular
- Language agnostic. Bril programs are just JSON.



BRIL is JSON

So are s-expressions

```
{  
  "functions": [  
    {  
      "name": "add5",  
      "args": [{"name": "n", "type": "int"}],  
      "type": "int",  
      "instrs": [  
        { "op": "const", "type": "int", "dest": "five", "value": 5 },  
        { "op": "add", "type": "int", "dest": "sum",  
          "args": ["n", "five"] },  
        { "op": "ret", "args": ["sum"] }  
      ]  
    },  
    {  
      "name": "main",  
      "args": [],  
      "instrs": [  
        { "op": "const", "type": "int", "dest": "a", "value": 9 },  
        { "op": "call", "type": "int", "dest": "b",  
          "funcs": ["add5"], "args": ["a"] },  
        { "op": "print", "args": ["b"] }  
      ]  
    }  
  ]  
}
```

```
(bril-define ((add5 int) (n int))  
  (set (five int) (const 5))  
  (set (sum int) (add n five))  
  (ret sum))
```

```
(bril-define ((main int))  
  (set (a int) (const 9))  
  (set (b int) (call add5 a))  
  (ret b))
```

```
[["bril-define", [["add5", "int"], ["n", "int"]],  
  ["set", ["five", "int"], ["const", 5]],  
  ["set", ["sum", "int"], ["add", "n", "five"]],  
  ["ret", "sum"]],  
 [["bril-define", [["main", "int"]],  
  ["set", ["a", "int"], ["const", 9]],  
  ["set", ["b", "int"], ["call", "add5", "a"]],  
  ["set", ["tmp", "int"],  
  ["call", "print", "b"]],  
  ["ret", "b"]]]]
```



Pass Around Programs, not Commands

- `guile dsl.scm | python llvm.py`
- No tight coupling between the different components beyond a JSON schema
- This JSON schema is literally a language
- Every language has got JSON parser, so can use whatever language we want
- Can replace parts easily



Layers of Languages

- **cat \${filename} | guile sexp-json.scm | python c-lisp.py | python brilisp.py | python llvm.py | bash run.sh**
- Simplify the translation by breaking it into layers which interact with each other via messages
- Messages are programs represented in text format

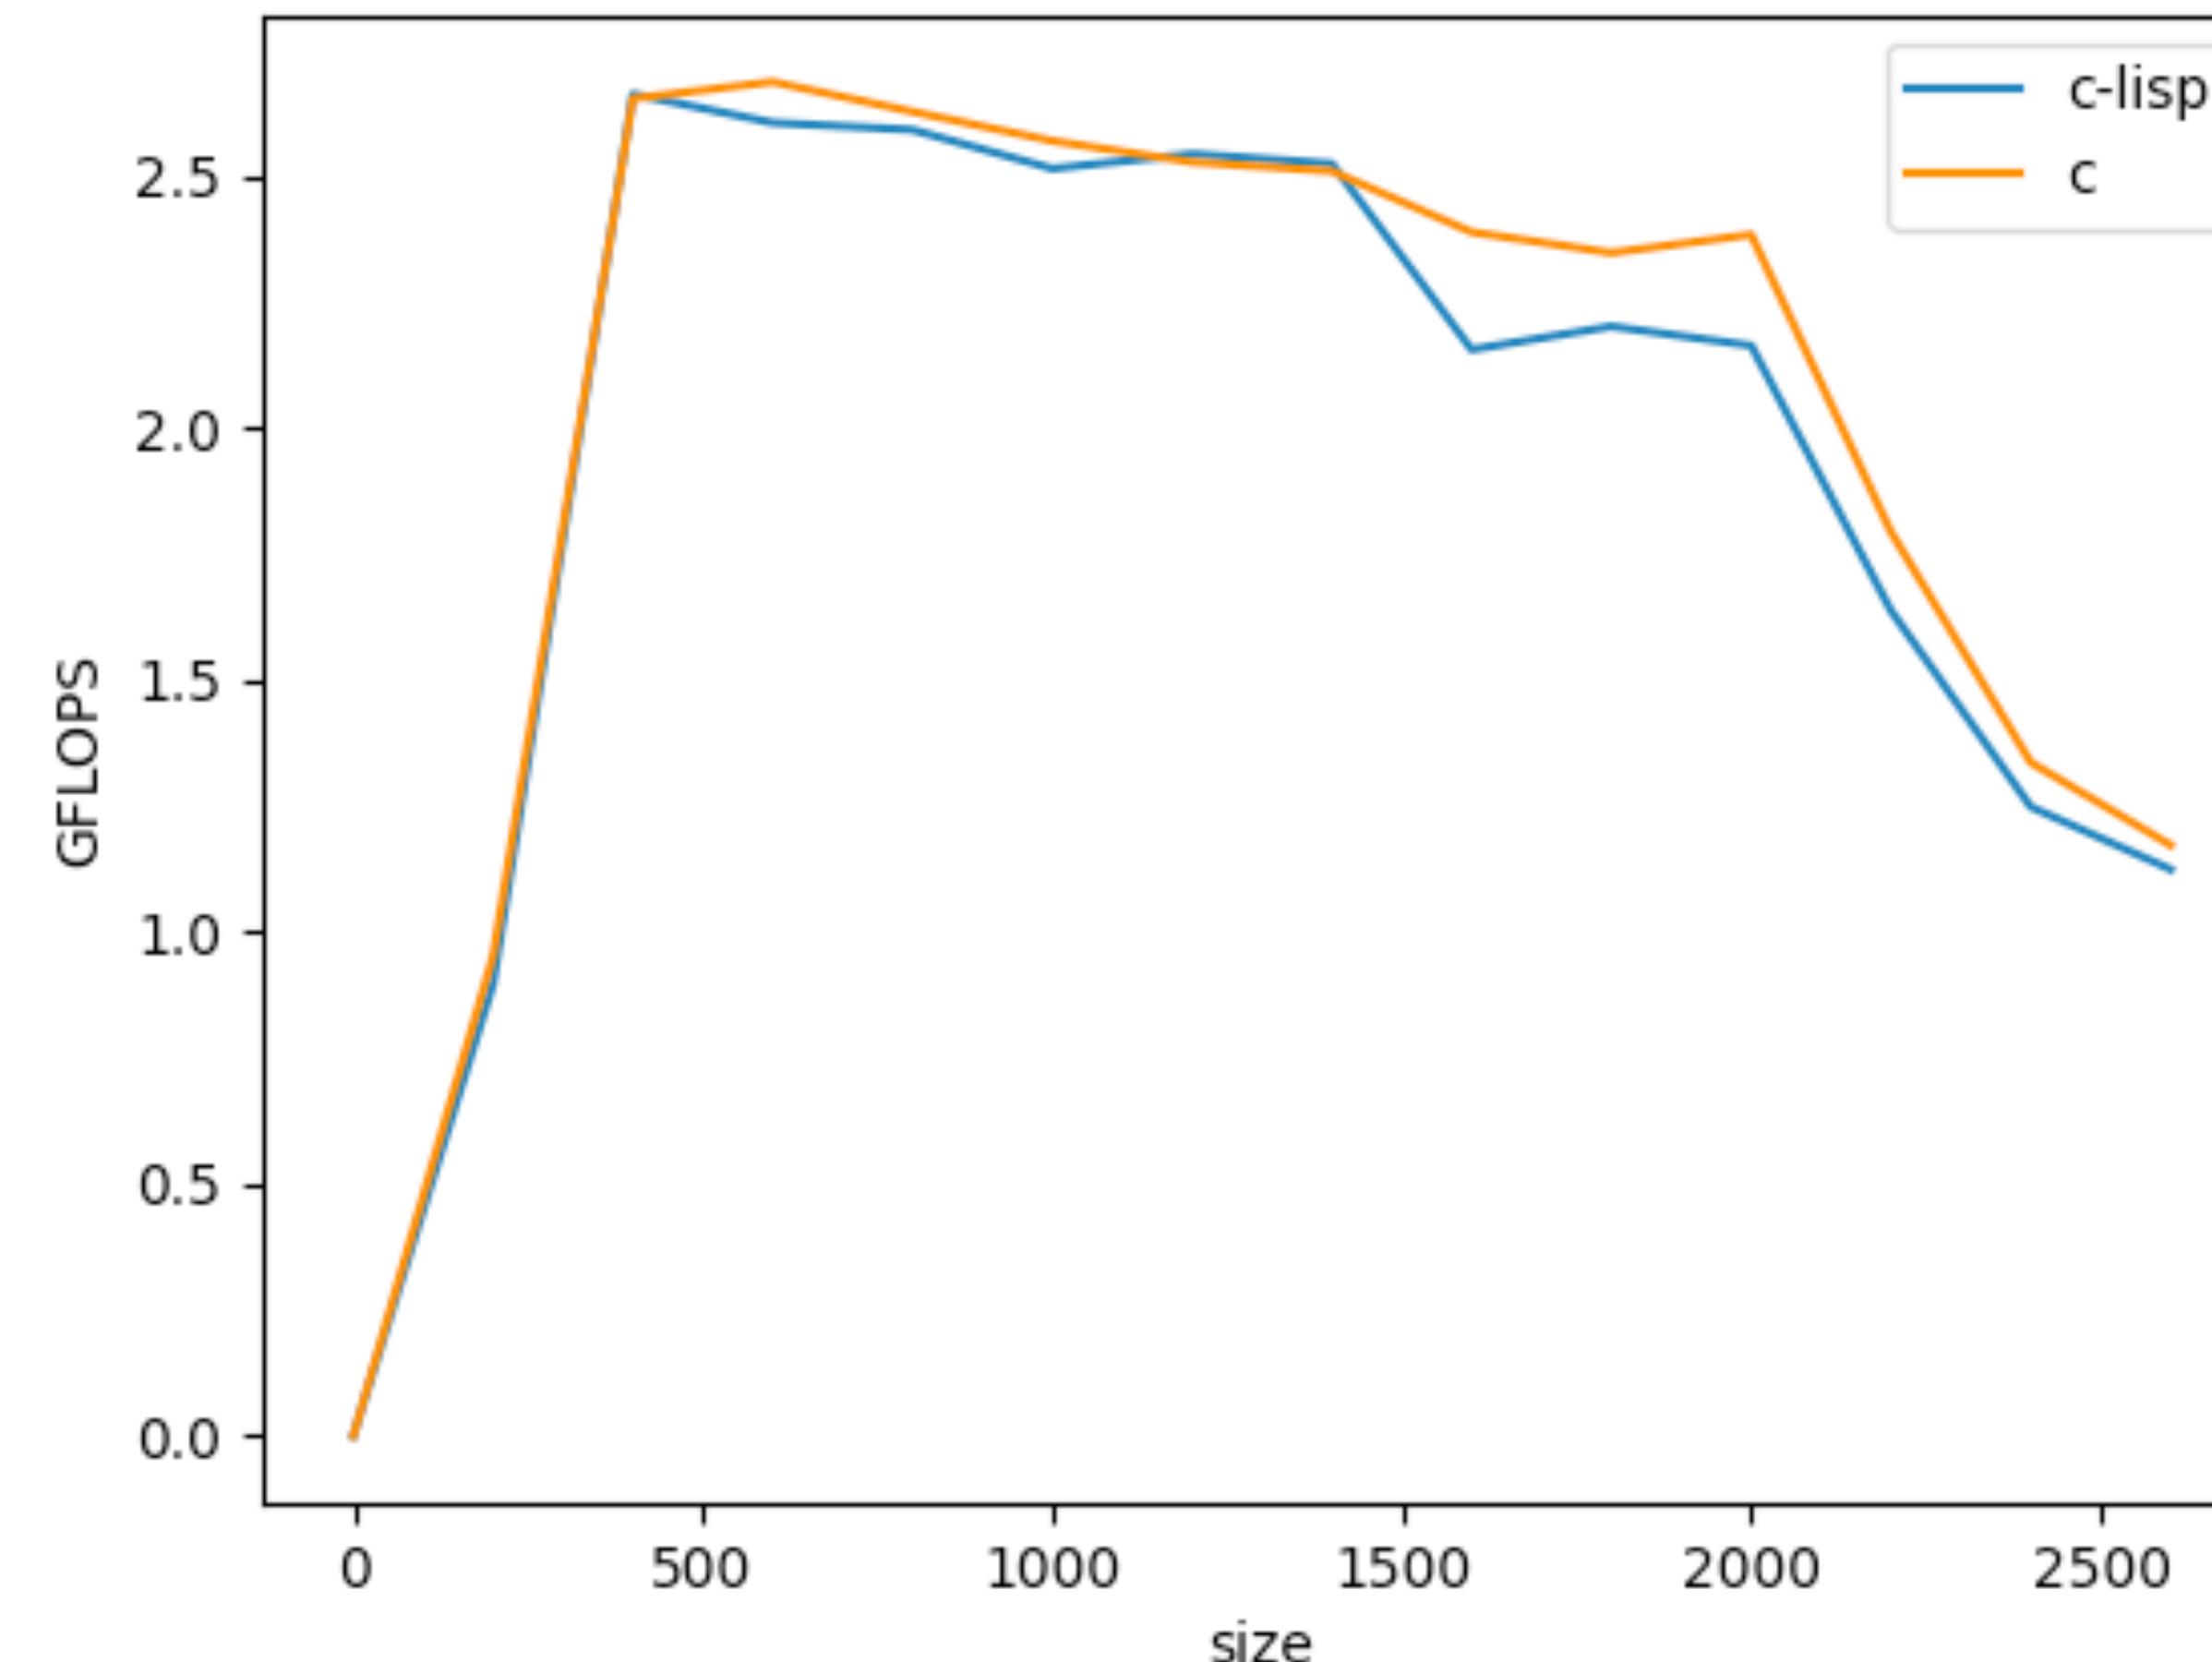


C-Lisp: C Like Frontend

So that we can generate this from other languages!

```
1 (define ((n-squares void) (n int))
2   ; Declare `i` as an integer
3   (declare (i int))
4   ; Iterate over the first n integers
5   (for ((set i 0)
6         (lt i n)
7         (set i (add i 1)))
8     ; `sq` is scoped to the `for` loop body
9     (declare (sq int))
10    (set sq (mul i i))
11    (call print sq))
12
13  (call print sq) ; Error: `sq` is undeclared
14  (ret))
```

Naive Matrix Multiplication Performance





GPU Support: WIP!

C-Lisp works almost as is!

chhasank / llama.lisp

Type ⌂ to search

ode Issues 20 Pull requests 7 Discussions Actions Projects Wiki Security Insights

Running C-Lisp code on Nvidia GPUs #71

Open GlowingScrewdr... wants to merge 1 commit into chhasank:main from GlowingScrewdriver:clisp-to-ptx

Conversation 0 Commits 1 Checks 0 Files changed 3

GlowingScrewdriver commented yesterday

This PR introduces a small framework for compiling and launching kernels written in C-Lisp on an Nvidia GPU

Bring up testing infra for kernels written in C-Lisp 09acf15

Add more commits by pushing to the [clisp-to-ptx](#) branch on [GlowingScrewdriver/llama.lisp](#).

This branch has not been deployed
No deployments

This branch has no conflicts with the base branch
Merging can be performed automatically.

Merge pull request You can also [open this in GitHub Desktop](#) or view [command line instructions](#).

Add a comment

```
(define ((kernel void) (a (ptr float)) (b (ptr float)) (c (ptr float)) (len int))
; Calculate c[i, j] from a[i, *] and b[*, j]
; a, b, c have dimensions len x len
(declare row int)
(declare col int)
(set row
      (add
        (call llvm.nvvm.read.ptx.sreg.tid.x)
        (mul (call llvm.nvvm.read.ptx.sreg.ntid.x) (call llvm.nvvm.read.ptx.sreg.ctaid.x)))
(set col
      (add
        (call llvm.nvvm.read.ptx.sreg.tid.y)
        (mul (call llvm.nvvm.read.ptx.sreg.ntid.y) (call llvm.nvvm.read.ptx.sreg.ctaid.y)))
(declare a-ptr (ptr float))
(declare b-ptr (ptr float))
(declare c-ptr (ptr float))
(set a-ptr (ptradd a (mul row len)))
(set b-ptr (ptradd b col))
(set c-ptr (ptradd c (add col (mul row len))))
(declare c-val float)
(set c-val 0.0)

(declare k int)
(for ((set k 0)
      (lt k len)
      (set k (add k 1)))
  (set c-val
      (fadd
        c-val
        (fmul (load a-ptr) (load b-ptr))))
  (set a-ptr (ptradd a-ptr 1))
  (set b-ptr (ptradd b-ptr len)))
(store c-ptr c-val)))
```



Future Work

- Custom GPU Driver
- Kernel generation
- Automatic program generation
- Advanced frontends like FL



Thanks!
Questions?
Please contribute!